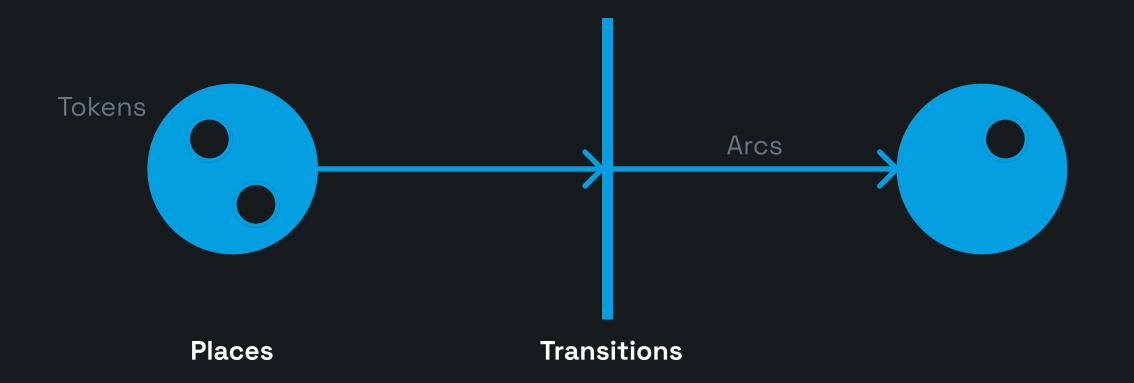
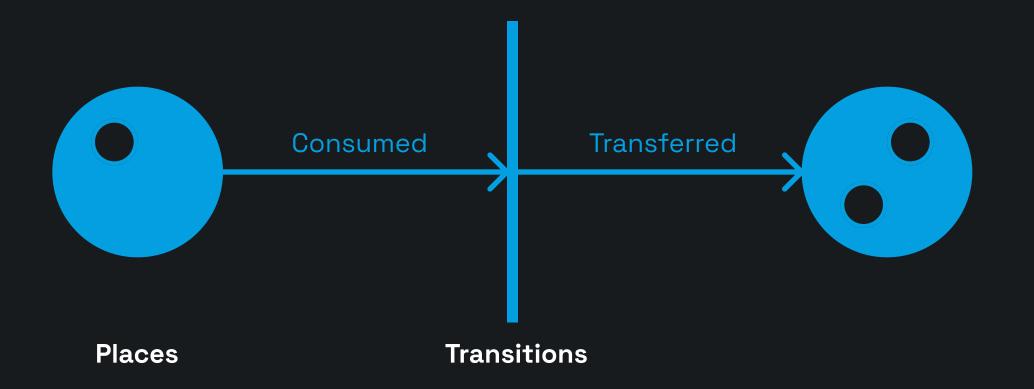
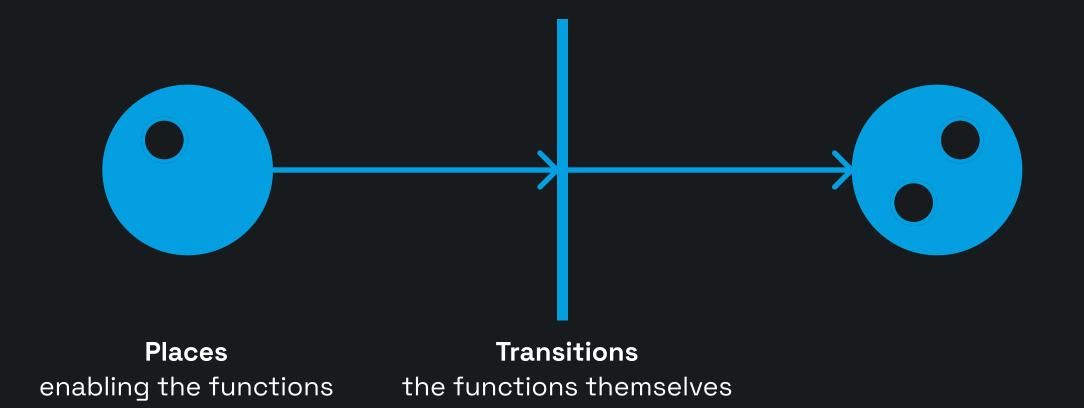
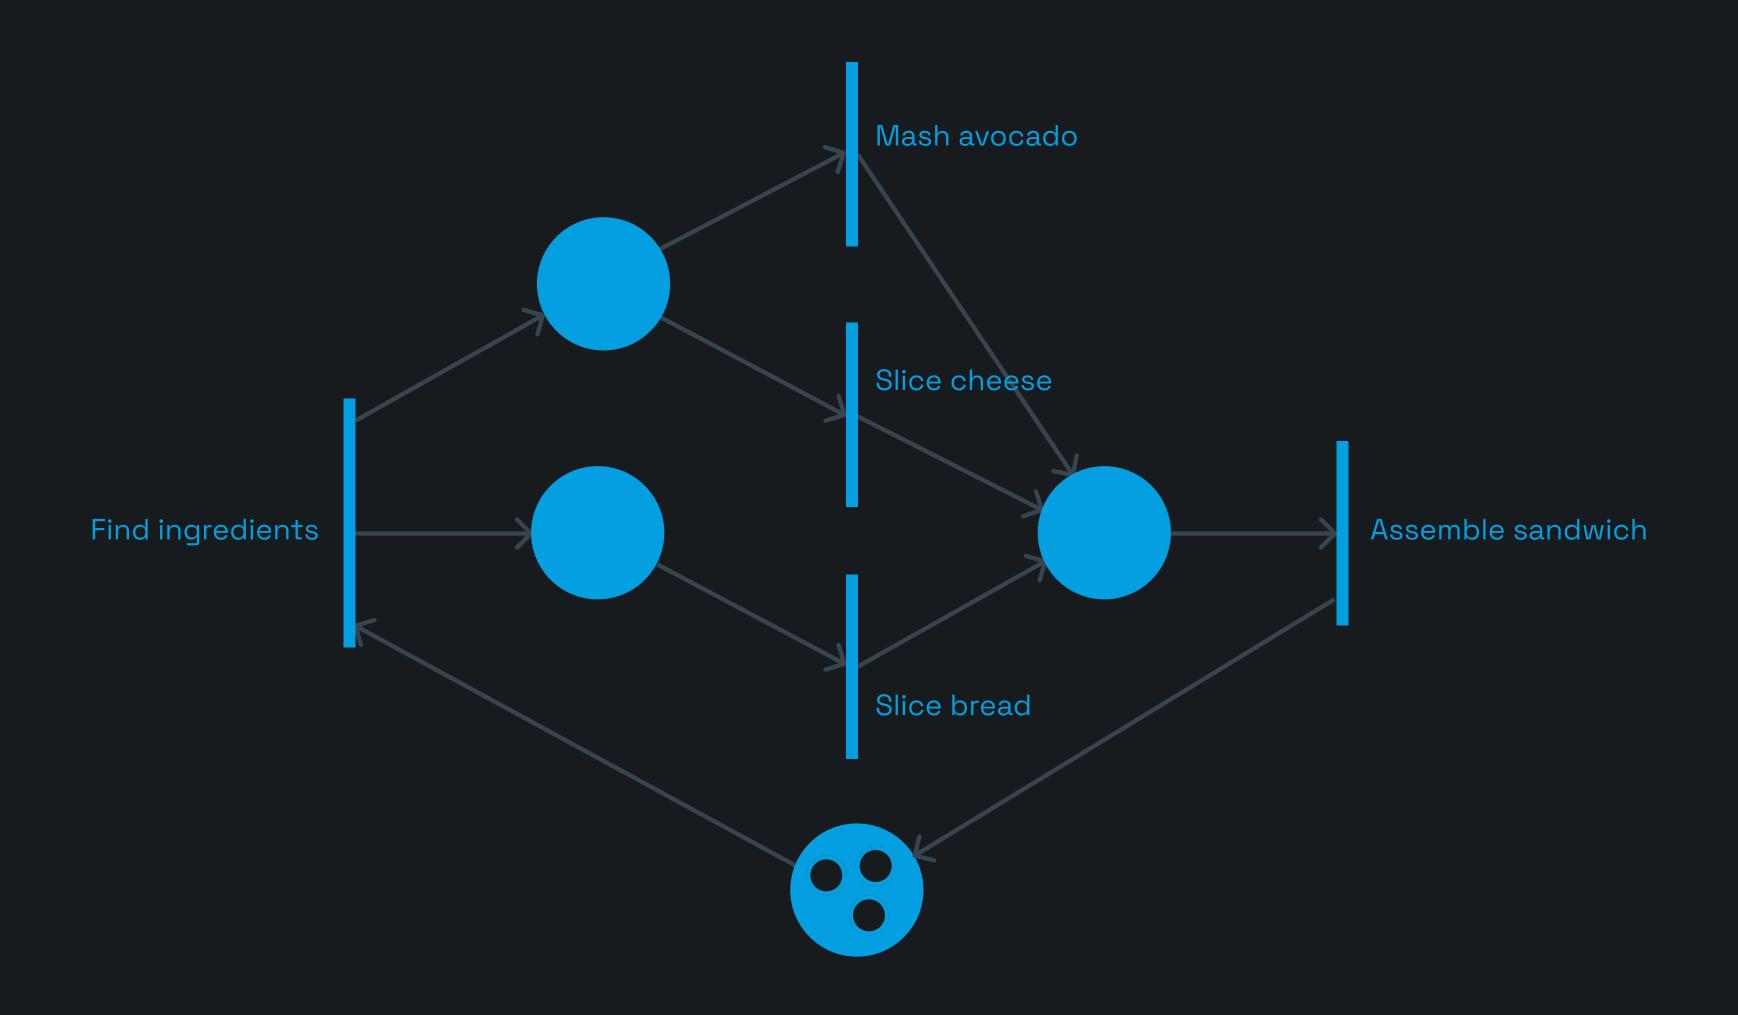
Fighting deadlocks with (modified) petri nets

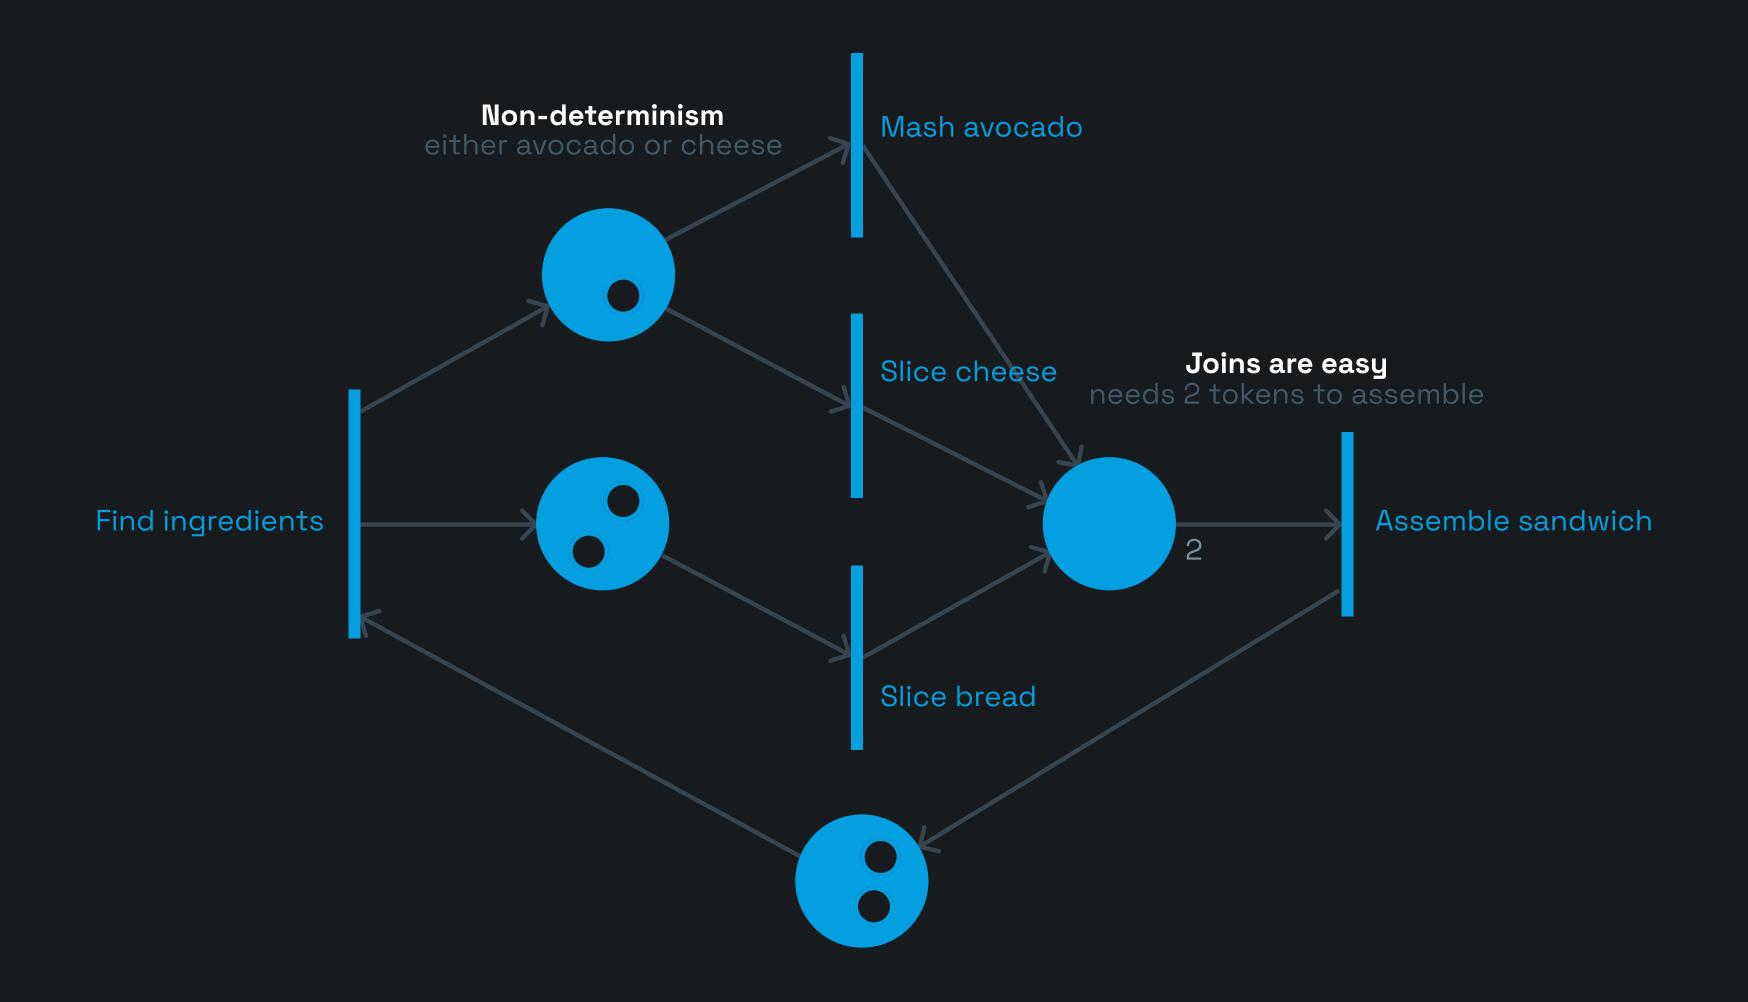




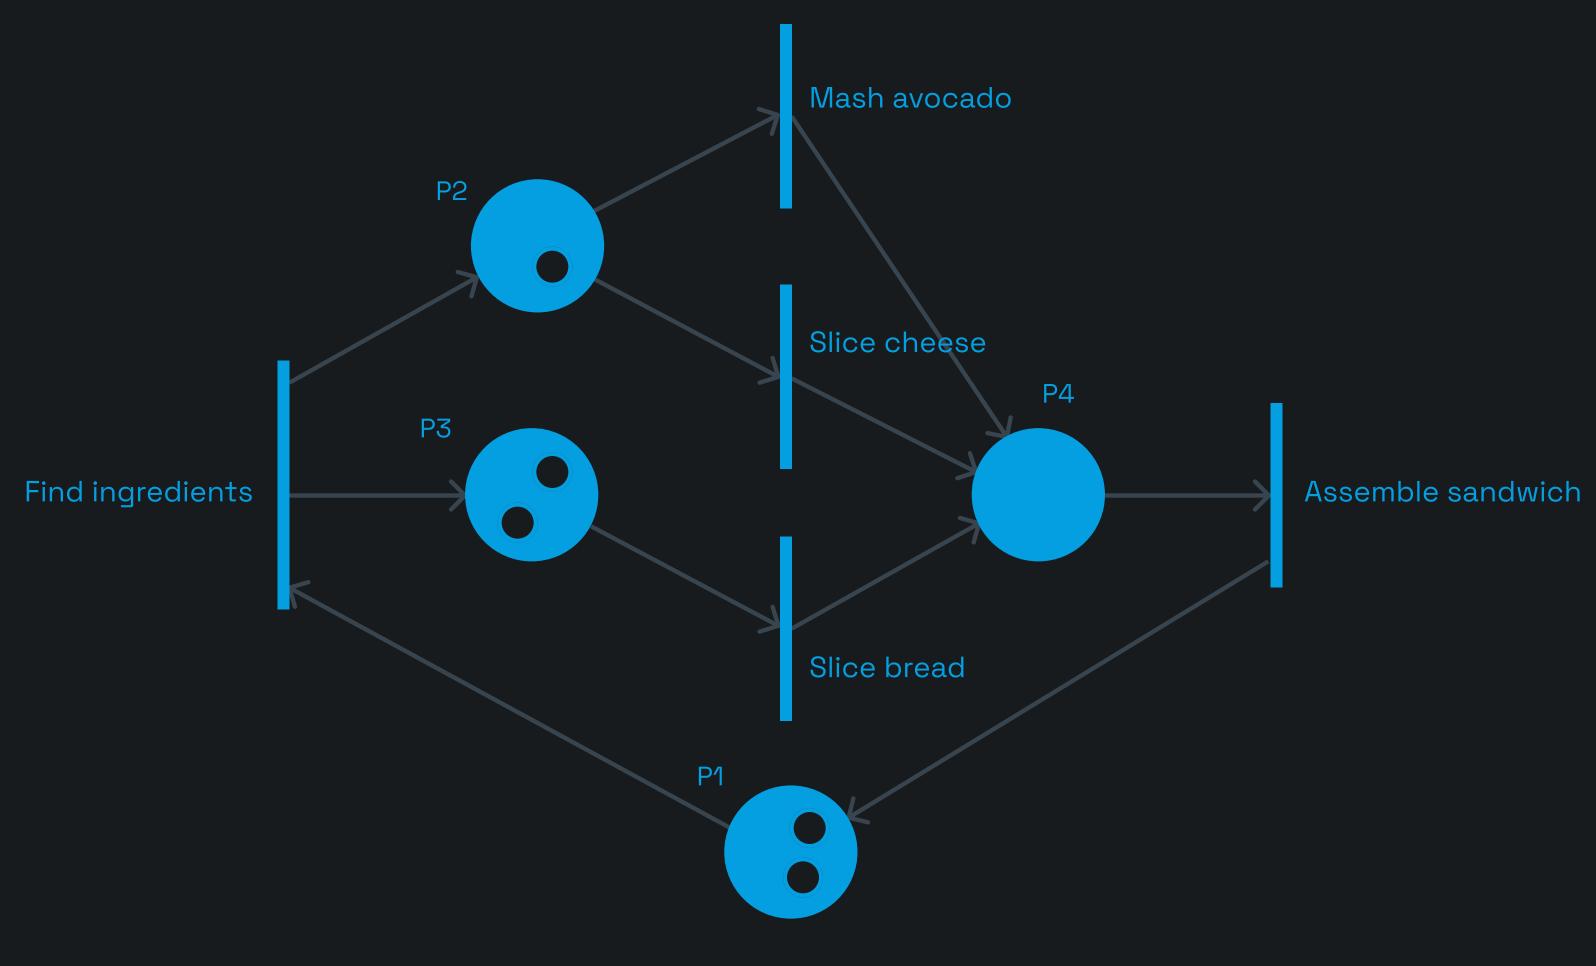
A software perspective





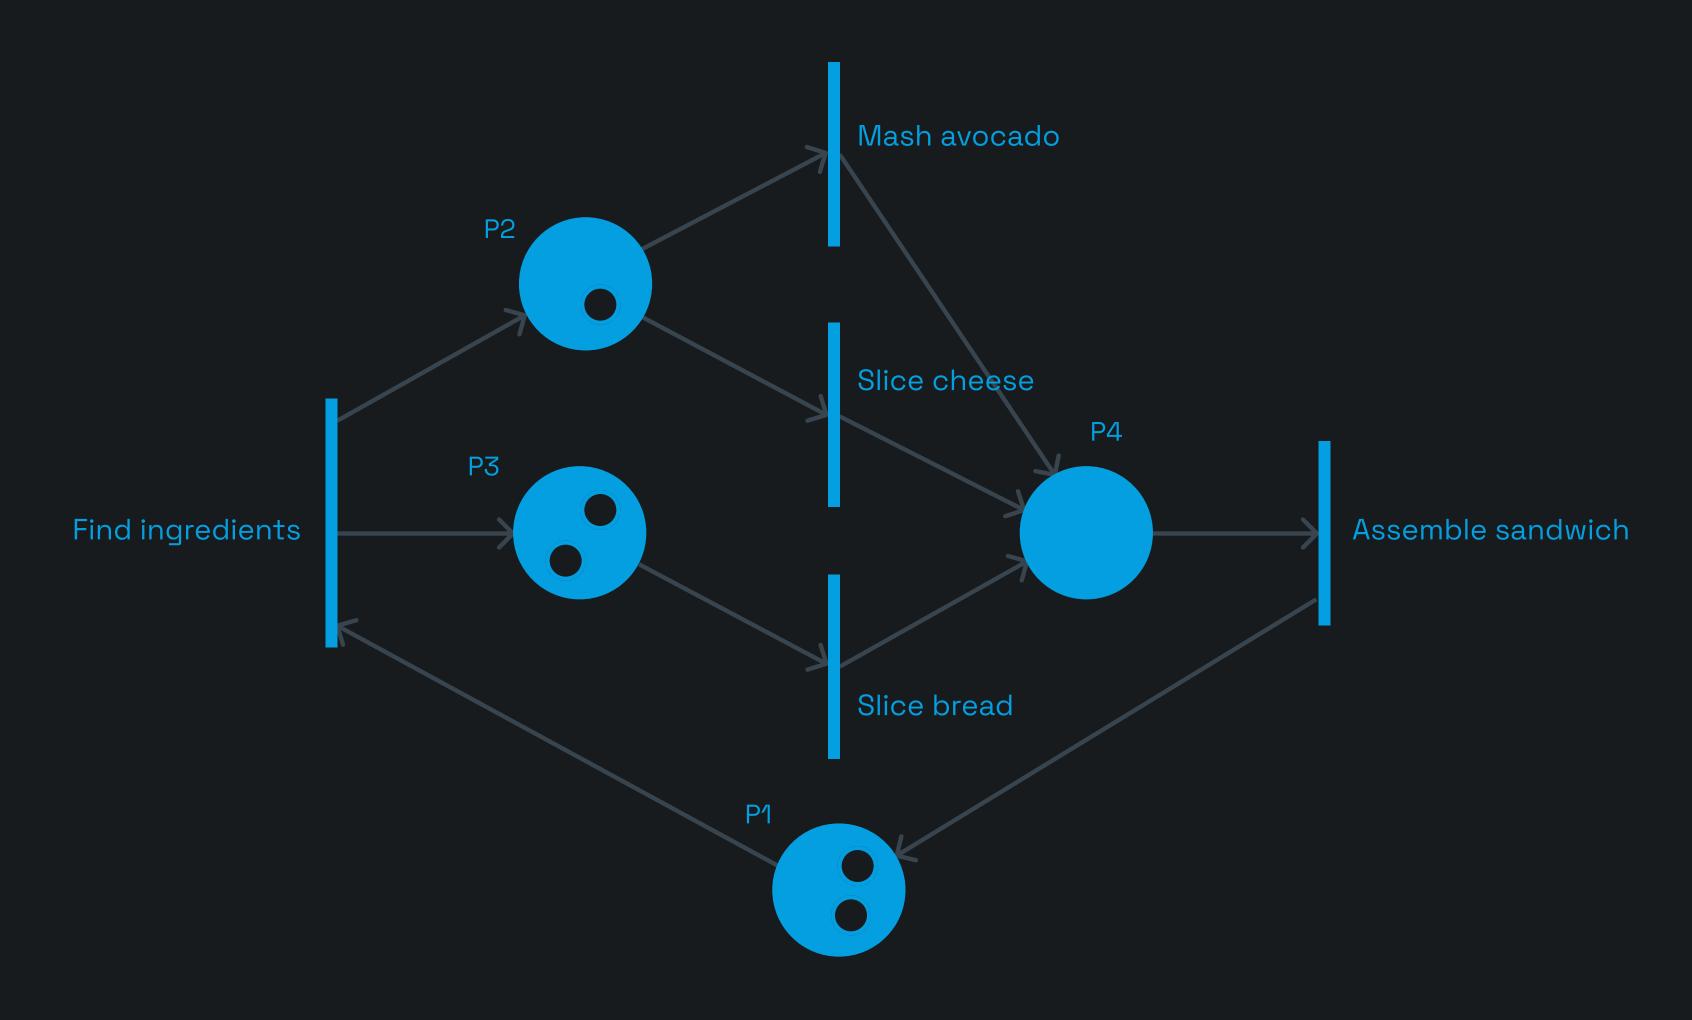


State (or marking)



M = [2, 1, 2, 0]

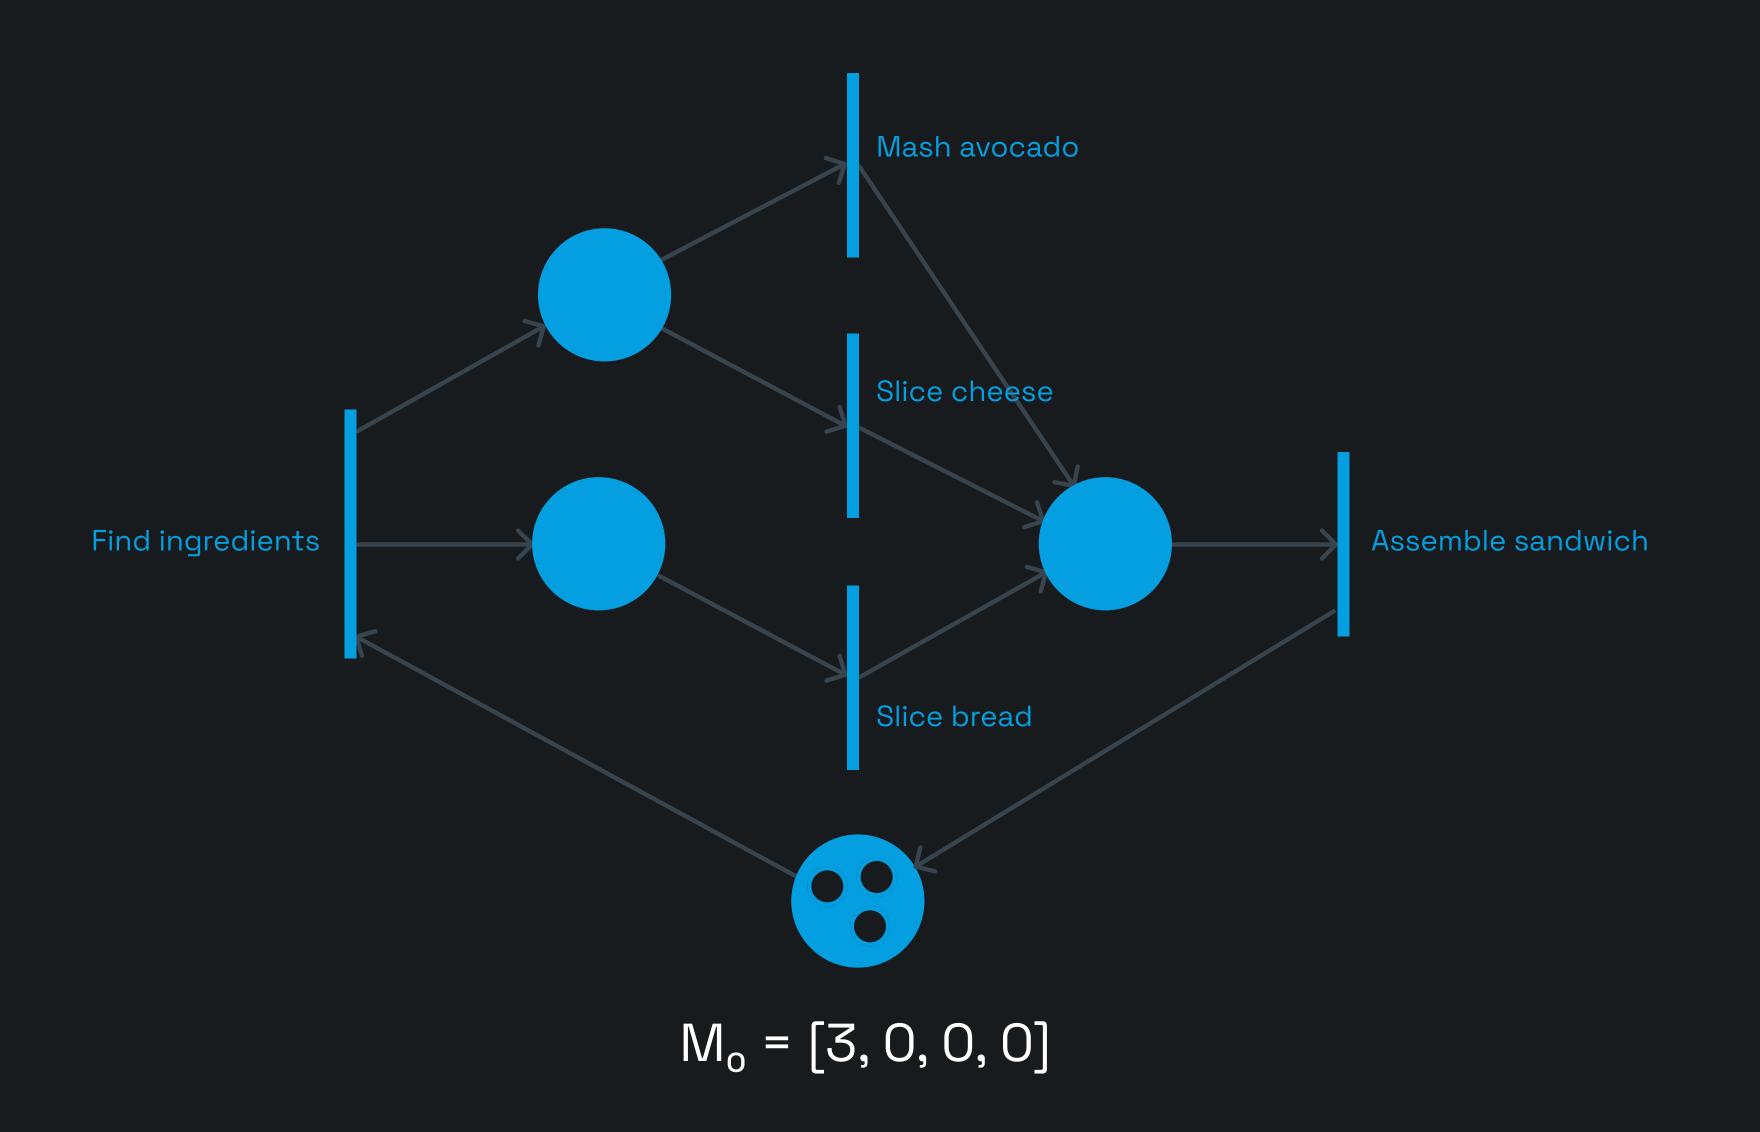
number of tokens in each place is the state of the net

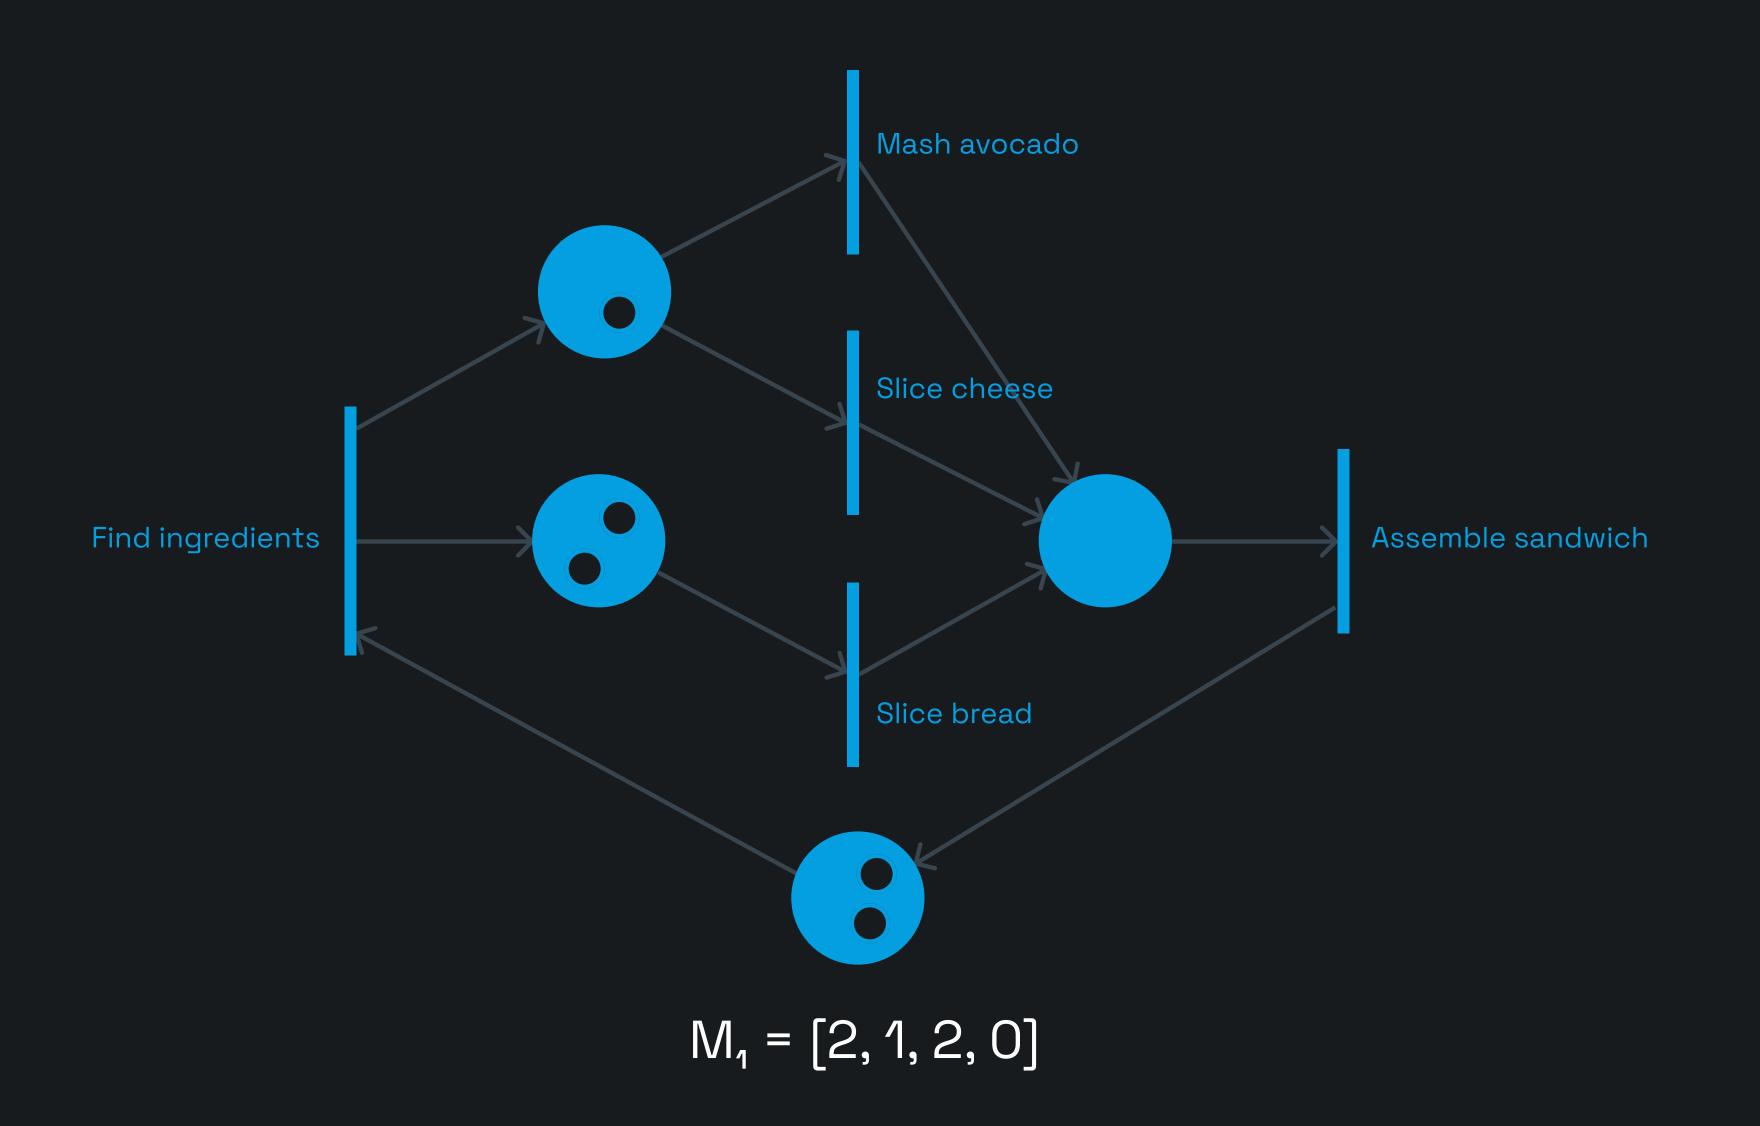


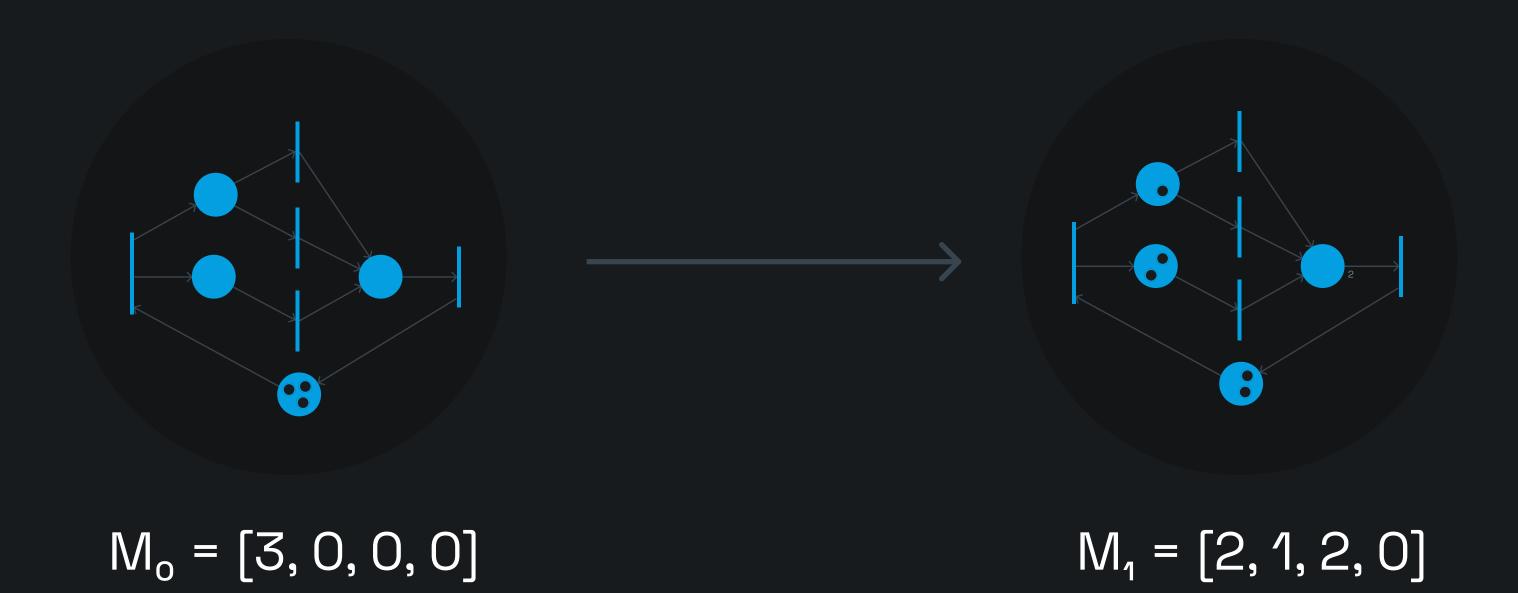
 $N = (P, T, A, W, M_o)$

Detecting deadlocks

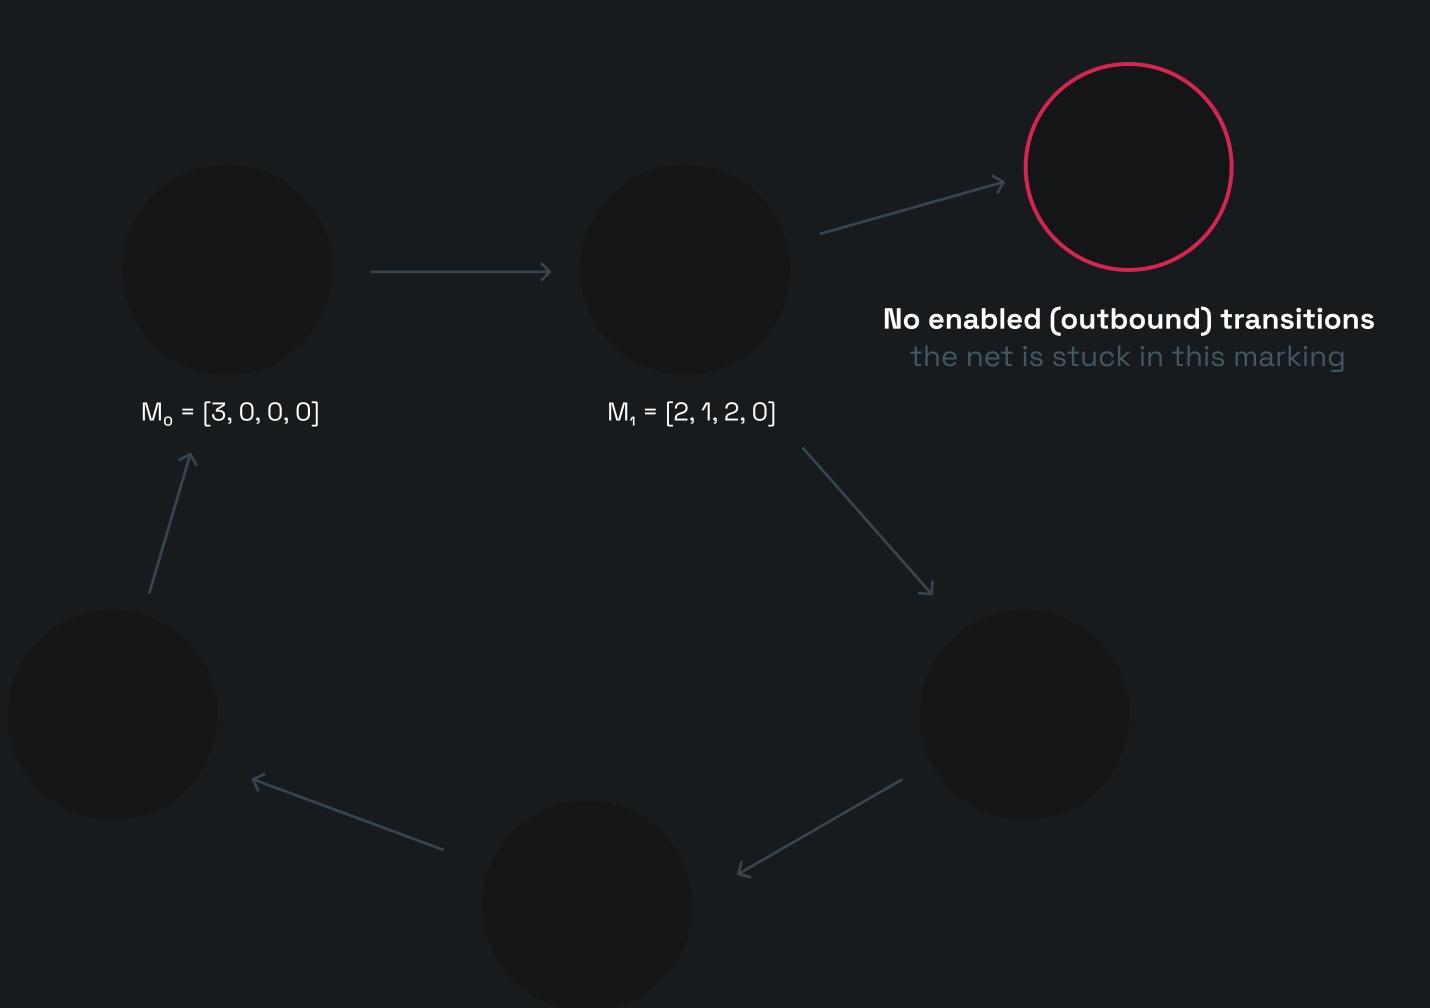
from just the initial marking Mo





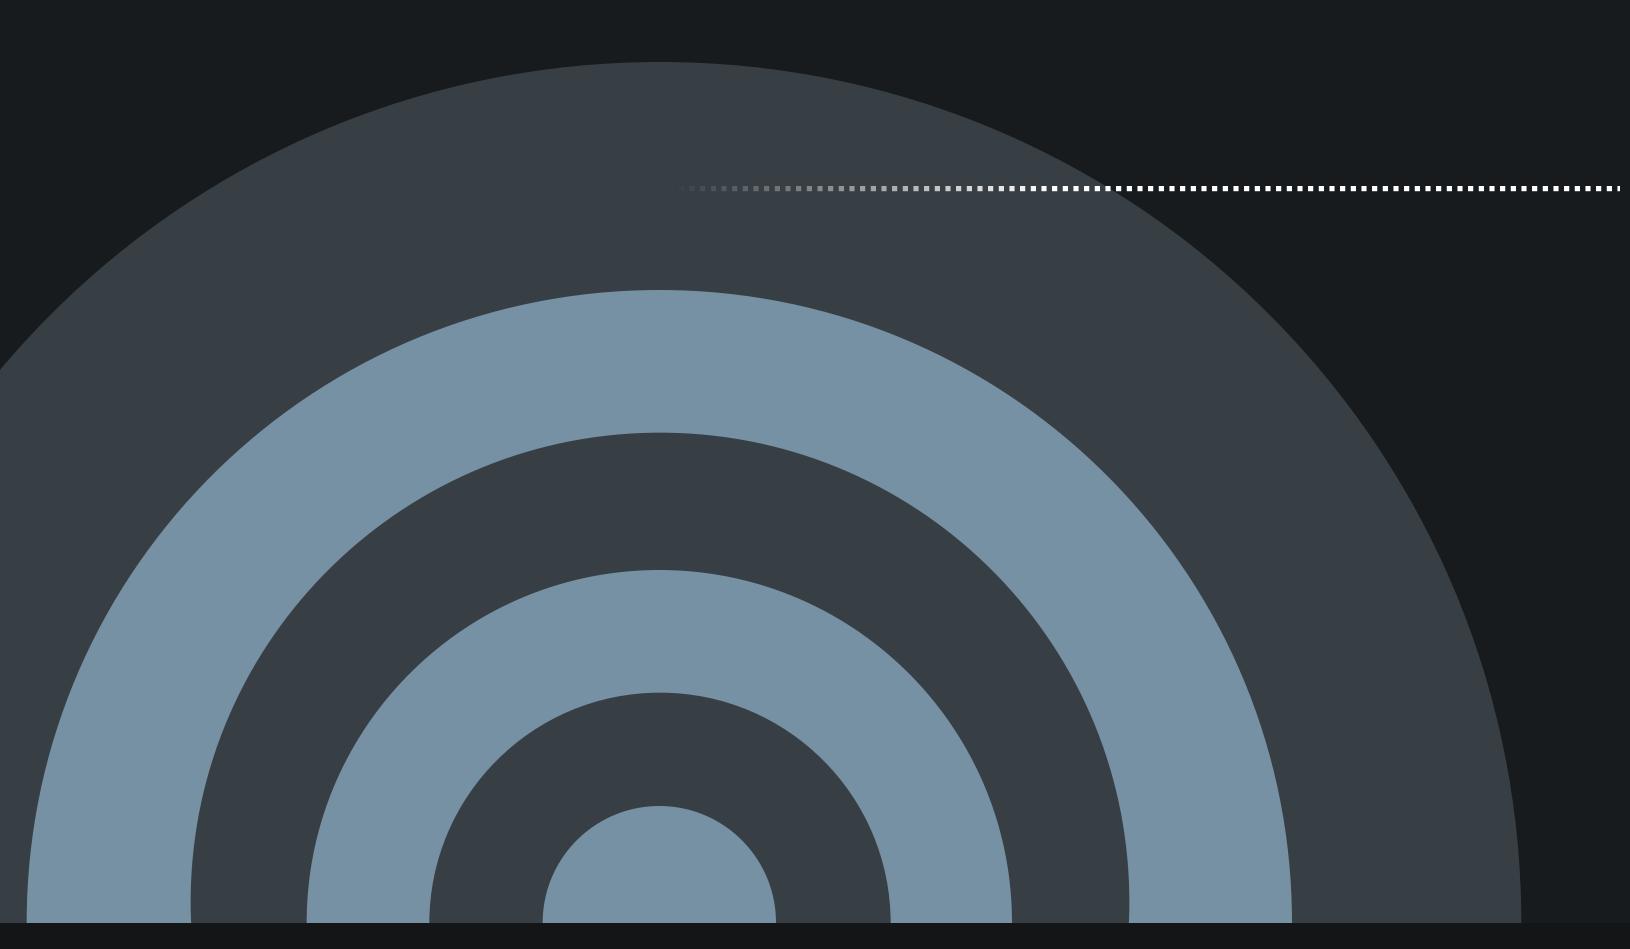






compile-time deadlock detection..

too good to be true?



Petri Net reachability analysis **EXPSPACE**

due to the **state explosion** problem

RESEARCH-ARTICLE

The Reachability Problem for Petri Nets Is Not Elementary

Authors: Wojciech Czerwiński, Sławomir Lasota, Ranko Lazić, Sźrôme Leroux, and Filip Mazowiecki Authors Info & Claims

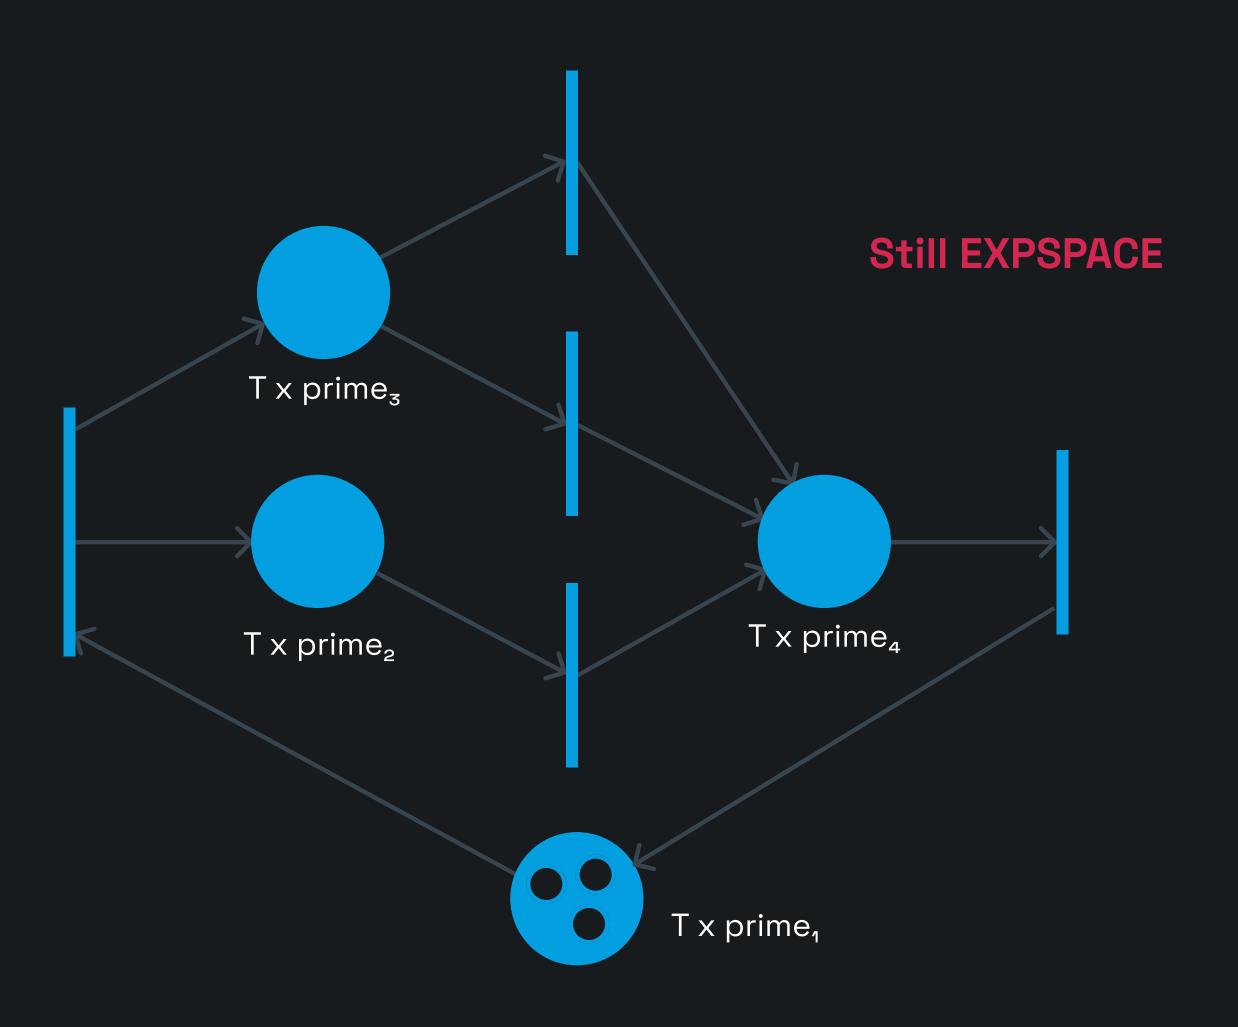
Journal of the ACM (JACM), Volume 68, Issue 1 • Article No.: 7, Pages 1 - 28 • https://doi.org/10.1145/3422822

Published: 22 December 2020 Publication History

of verification. Decidability was proved by Mayr in his seminal STOC 1981 work, and, currently, the best published upper bound is non-primitive recursive Ackermannian of Leroux and Schmitz from Symposium on Logic in Computer Science 2019. We establish a non-elementary lower bound, i.e., that the reachability problem needs a tower of exponentials of time and space. Until this work, the best lower bound has been exponential space, due to Lipton in 1976. The new lower bound is a major breakthrough for several reasons. Firstly, it shows that the reachability problem is much harder than the

Partial hashing using the FTA

hash and un-hash based on unique primes and places that changed



Z-Net

Modifying petri nets with shady math for polynomial-time deadlock detection

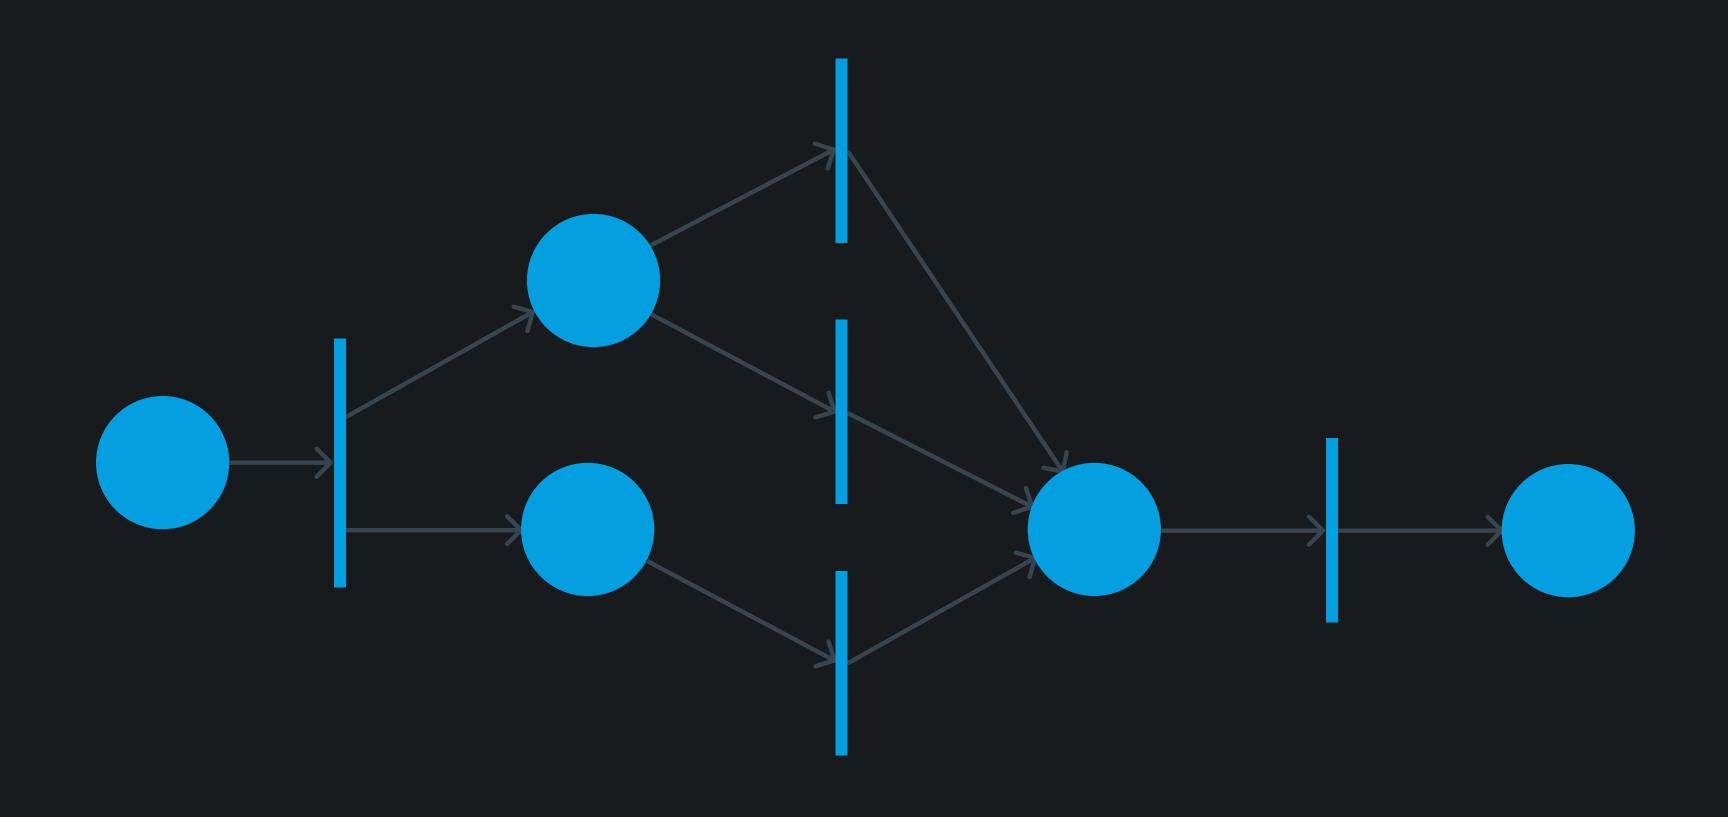
Goals

- Polynomial time deadlock detection
- Recursion* and free-choice

*causes state explosion by default

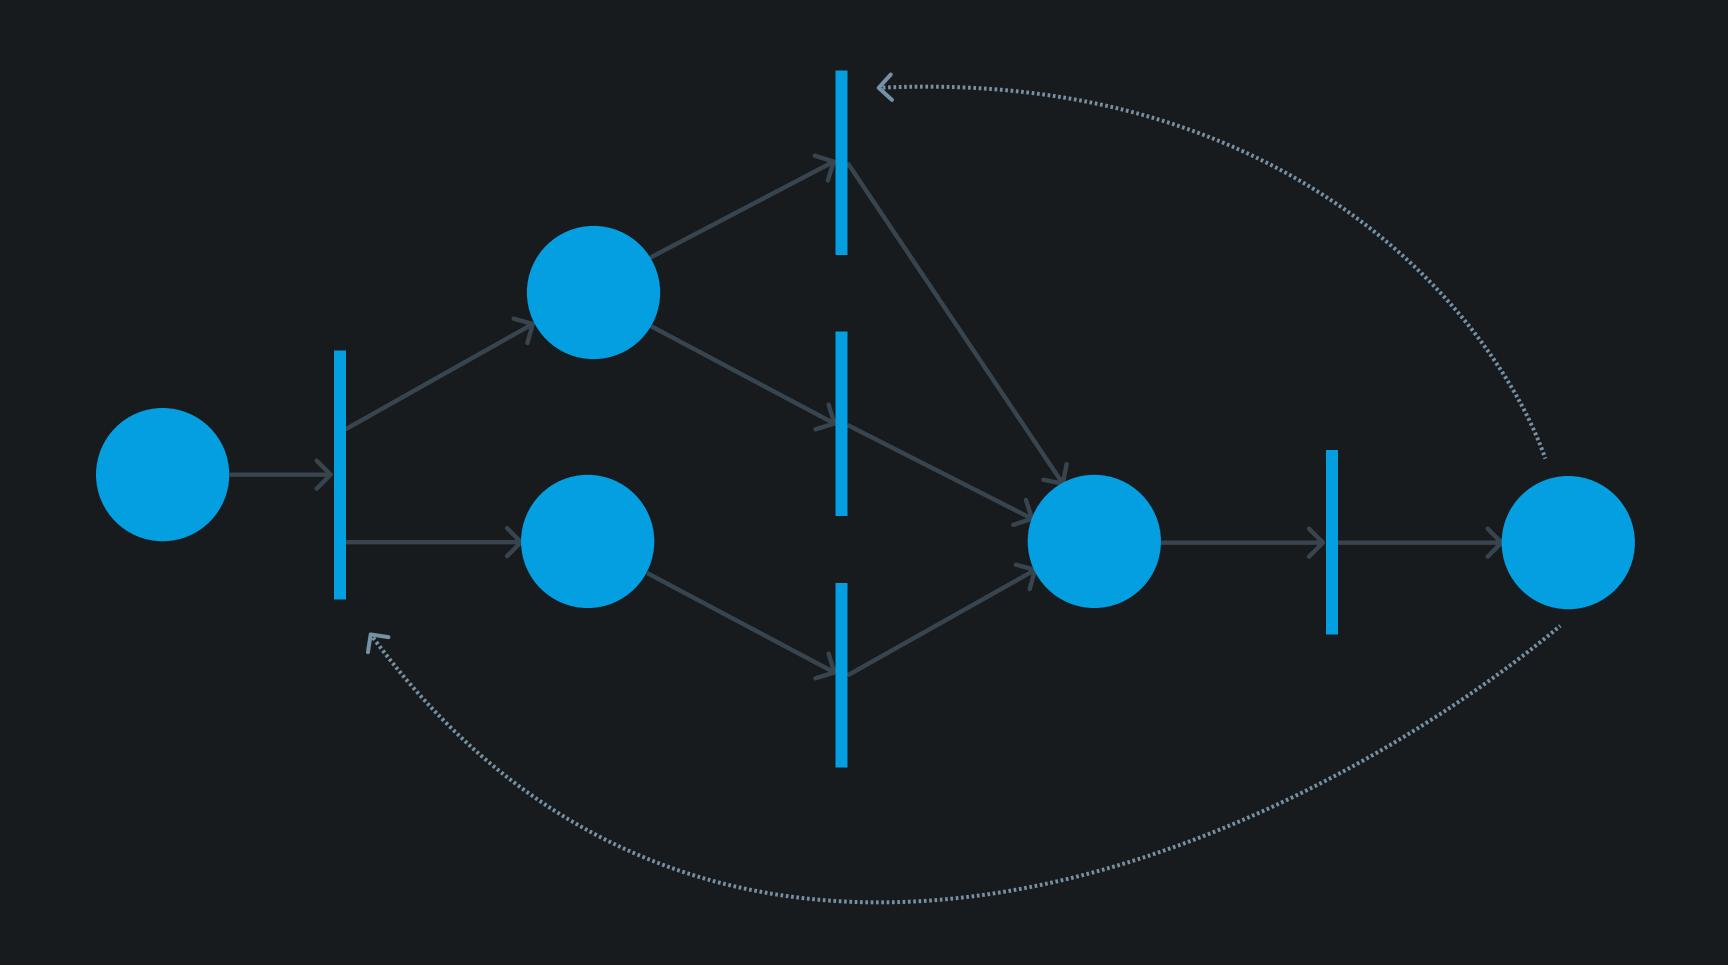
Reachability in acyclic petri nets is manageable

but how do we model recursion?

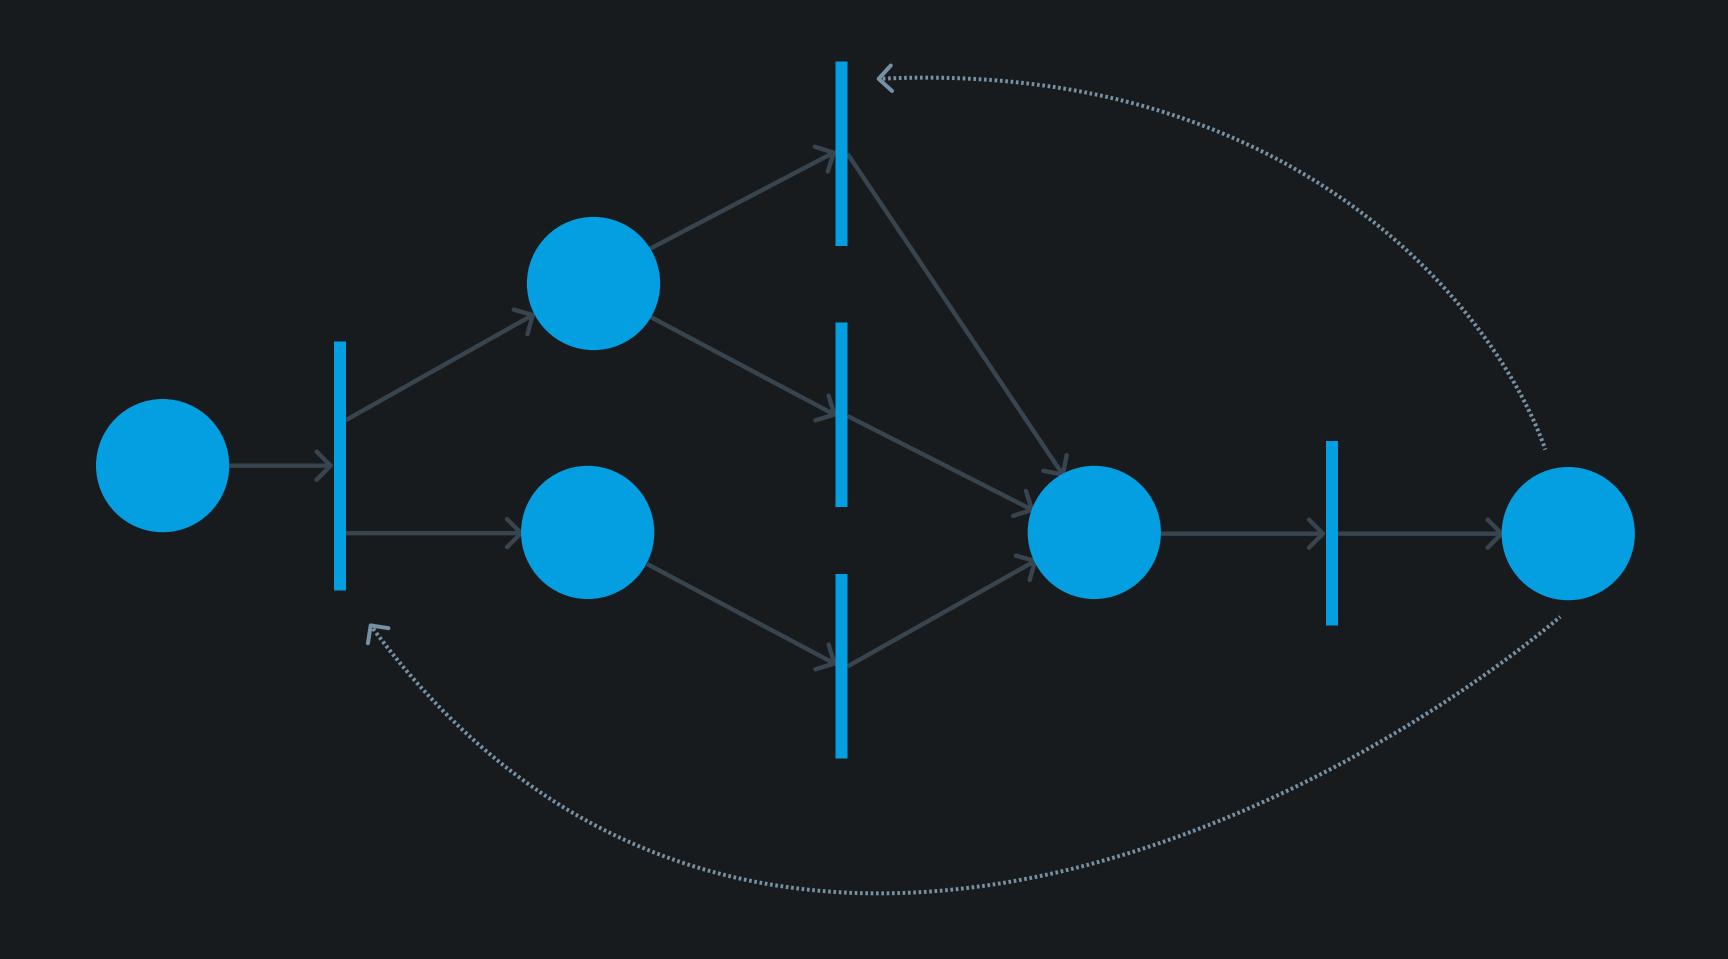


Reachability in acyclic petri nets is manageable

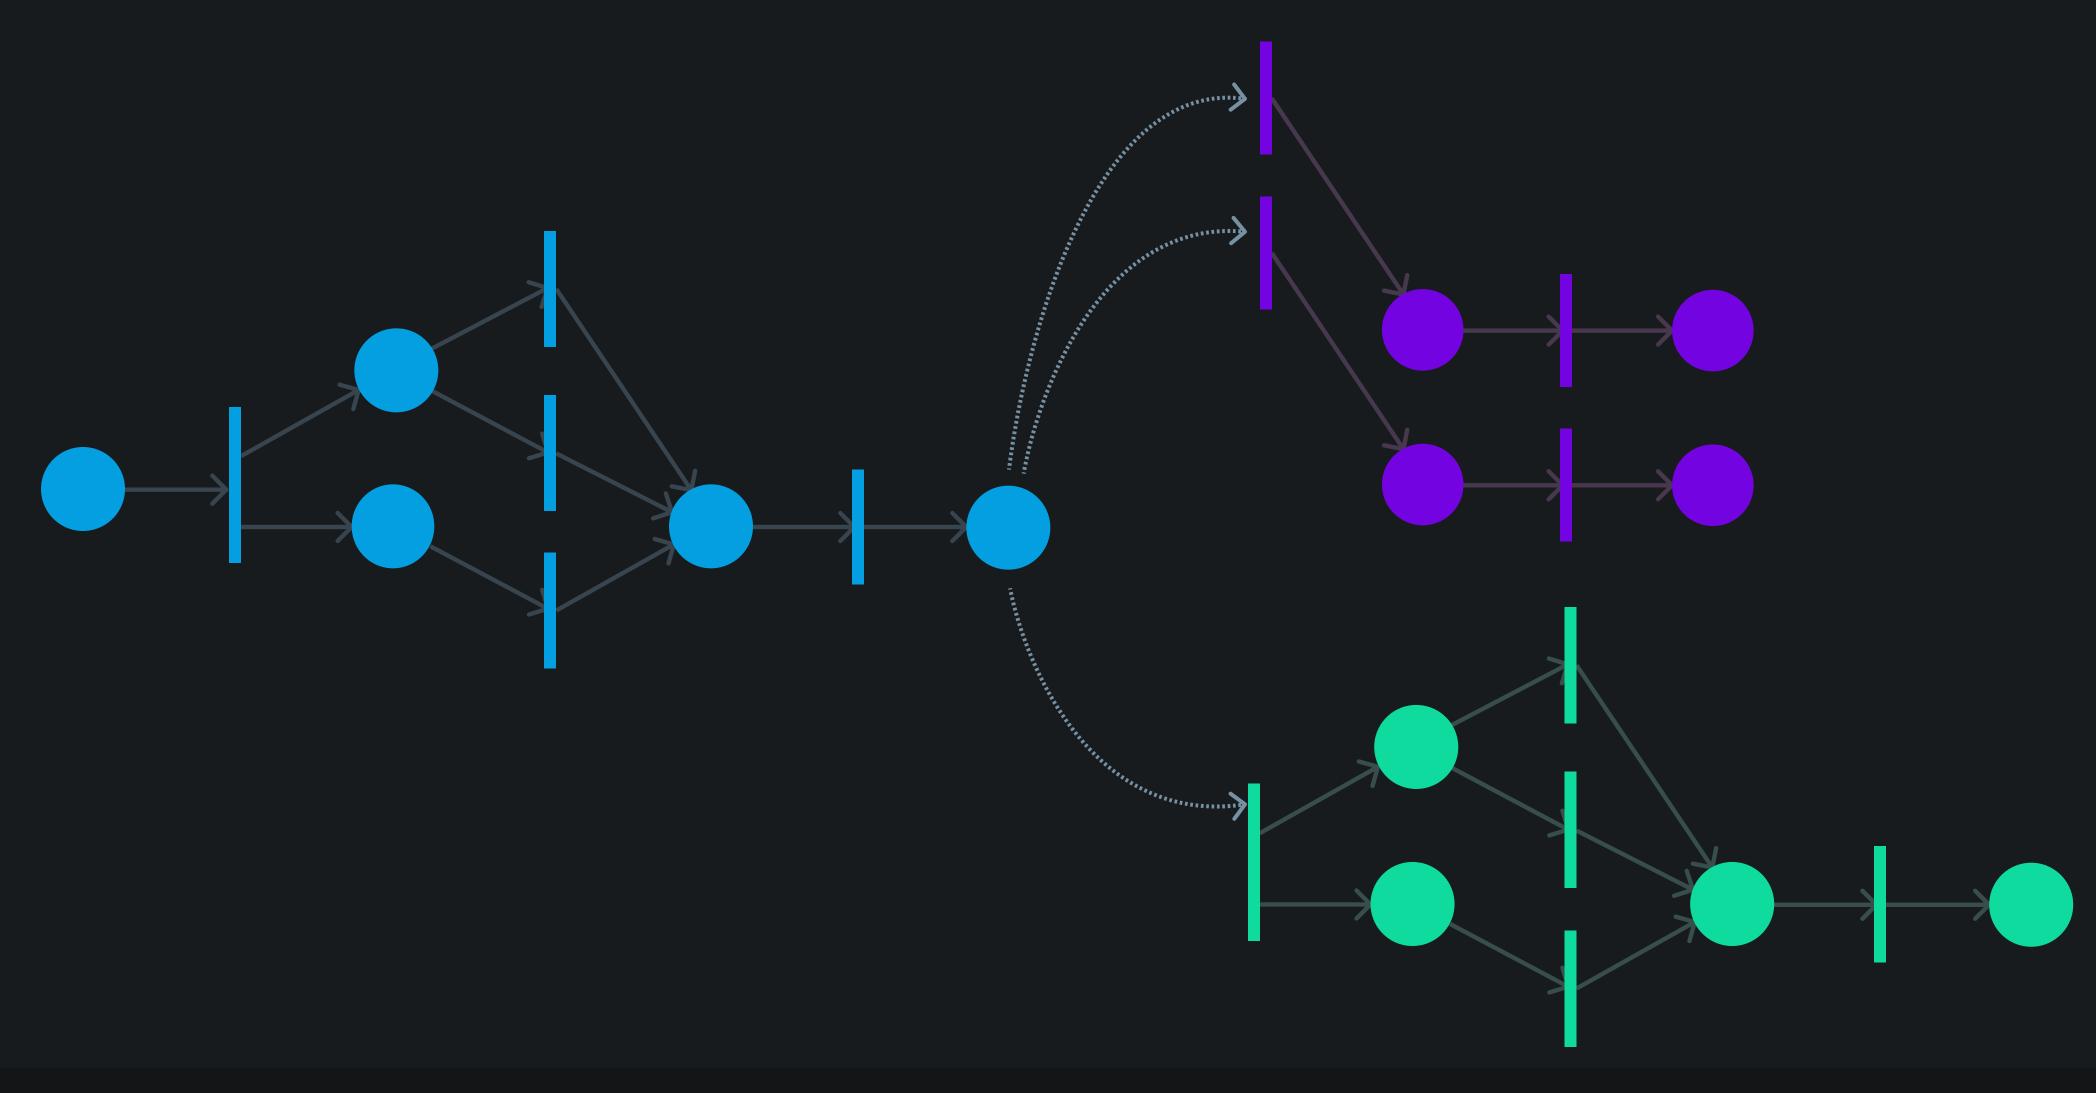
but how do we model recursion?



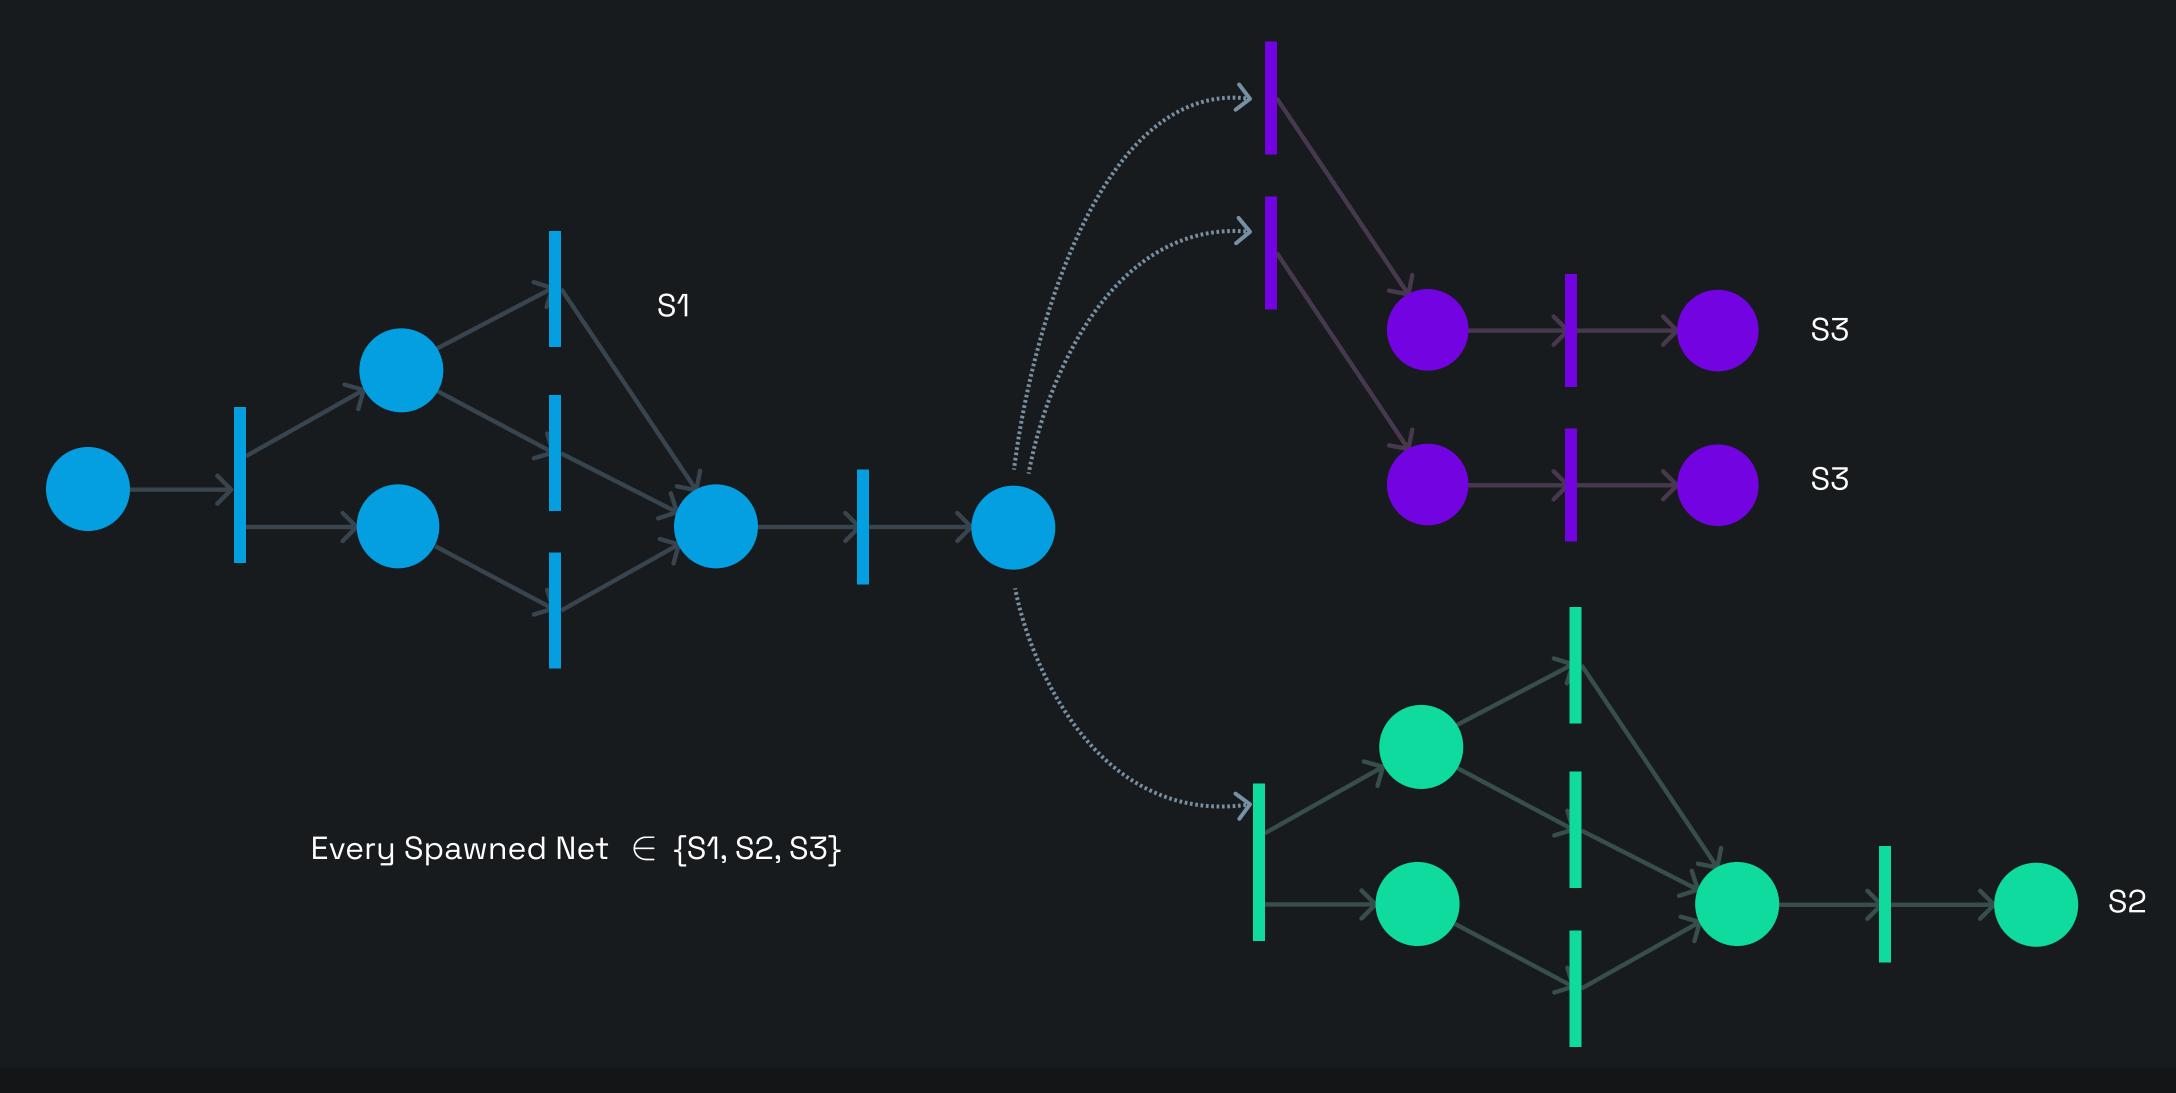
... by not using recursion, and spawning state-independent subnets

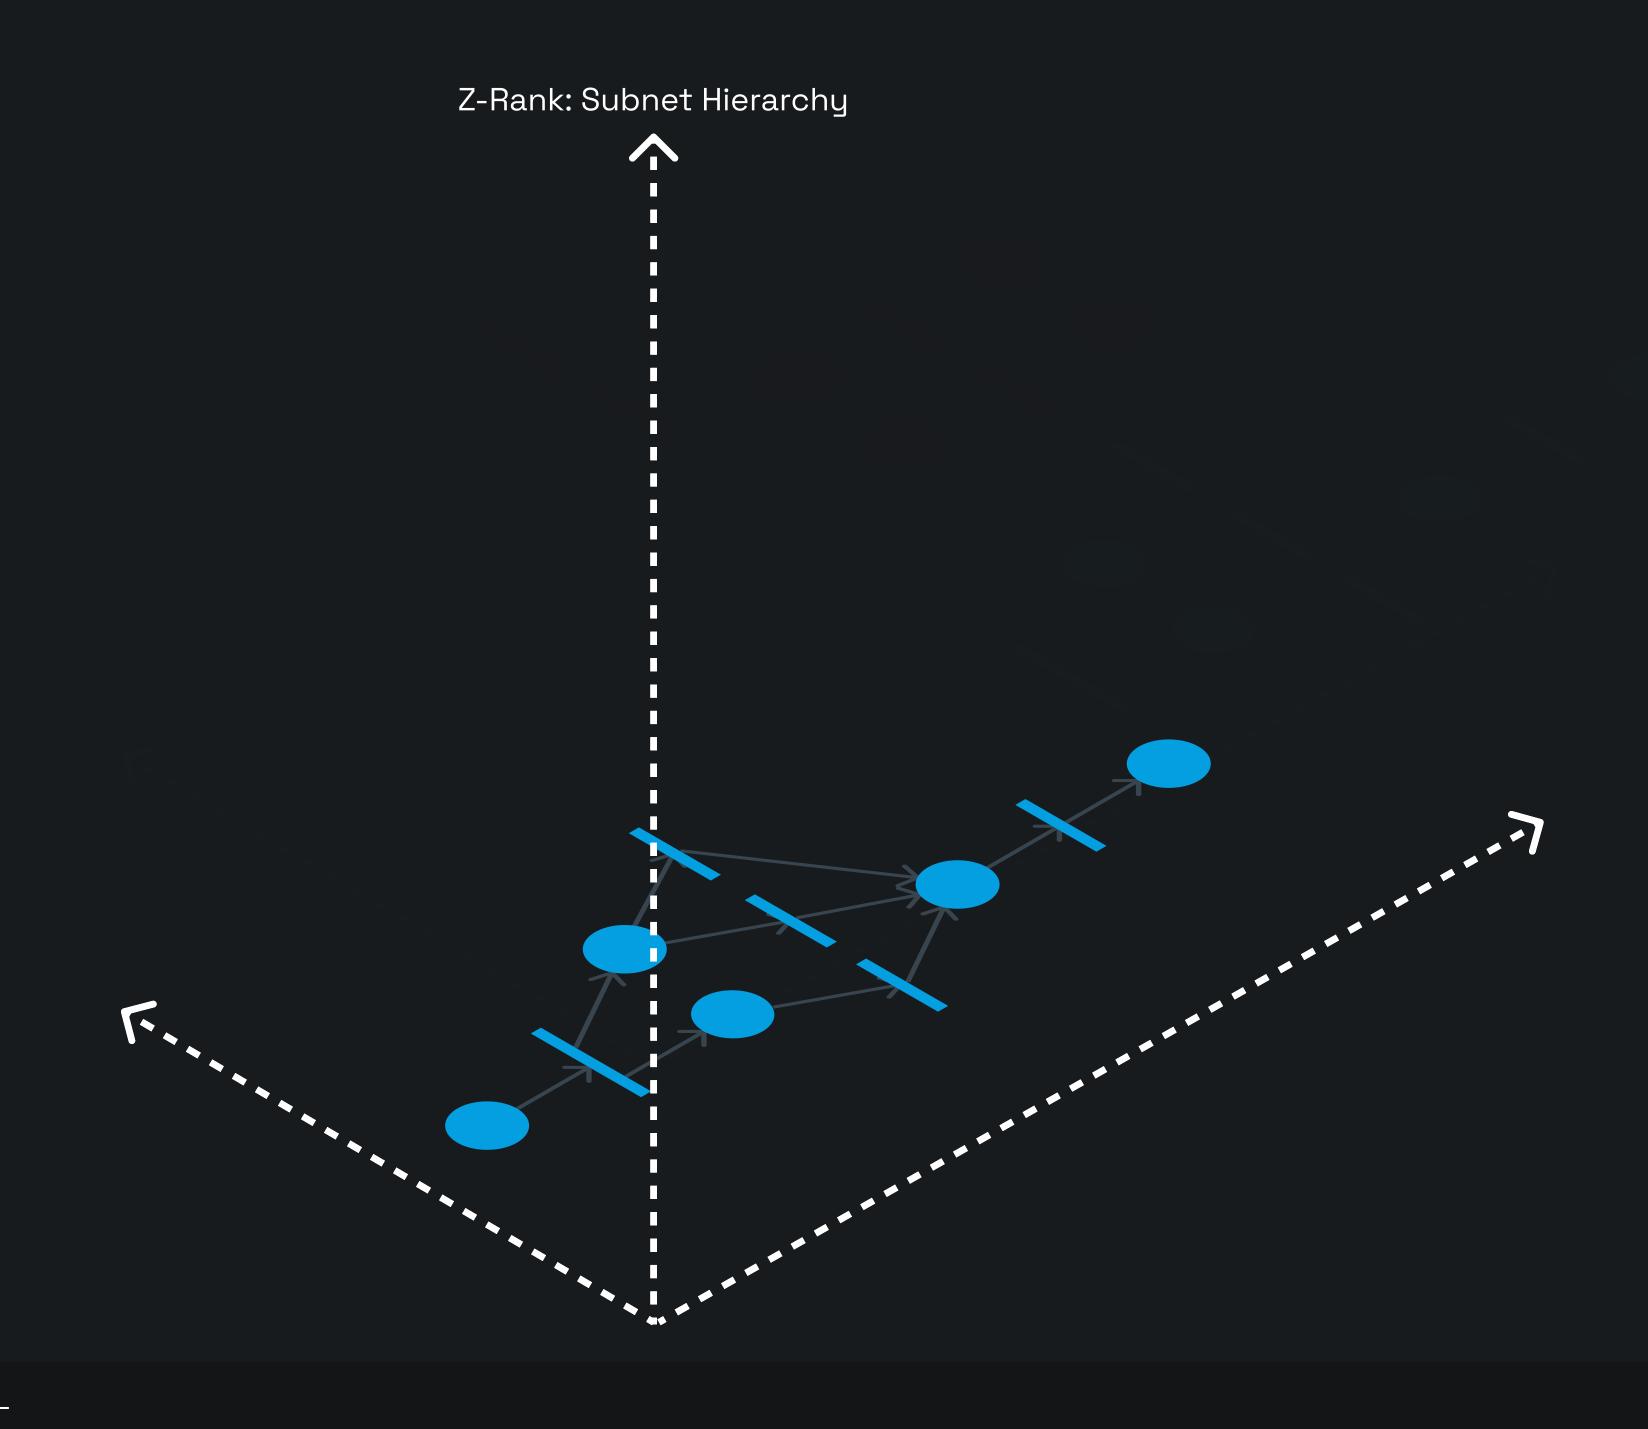


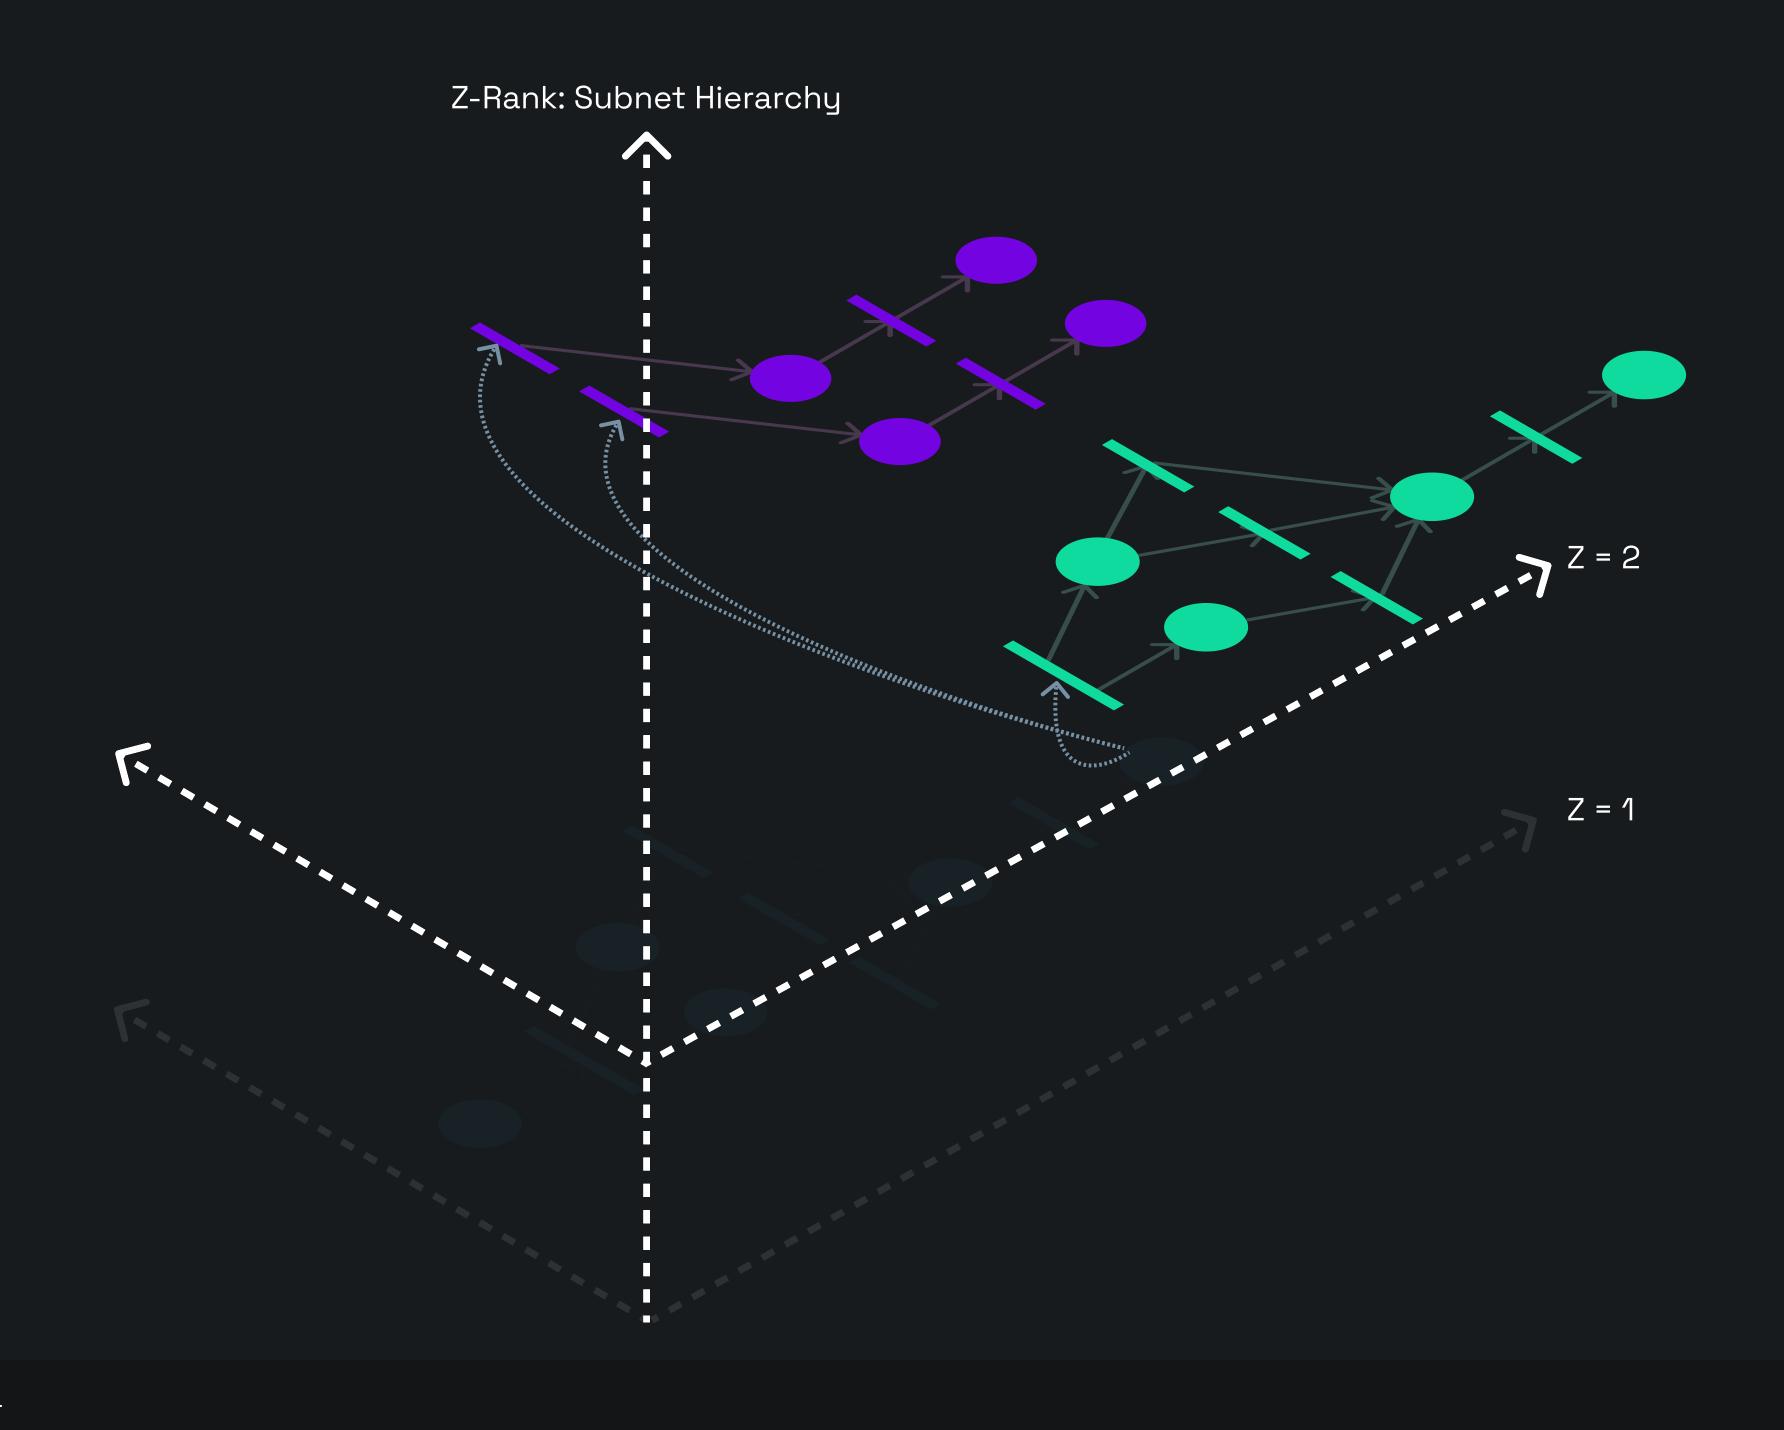
... by not using recursion, and spawning state-independent subnets

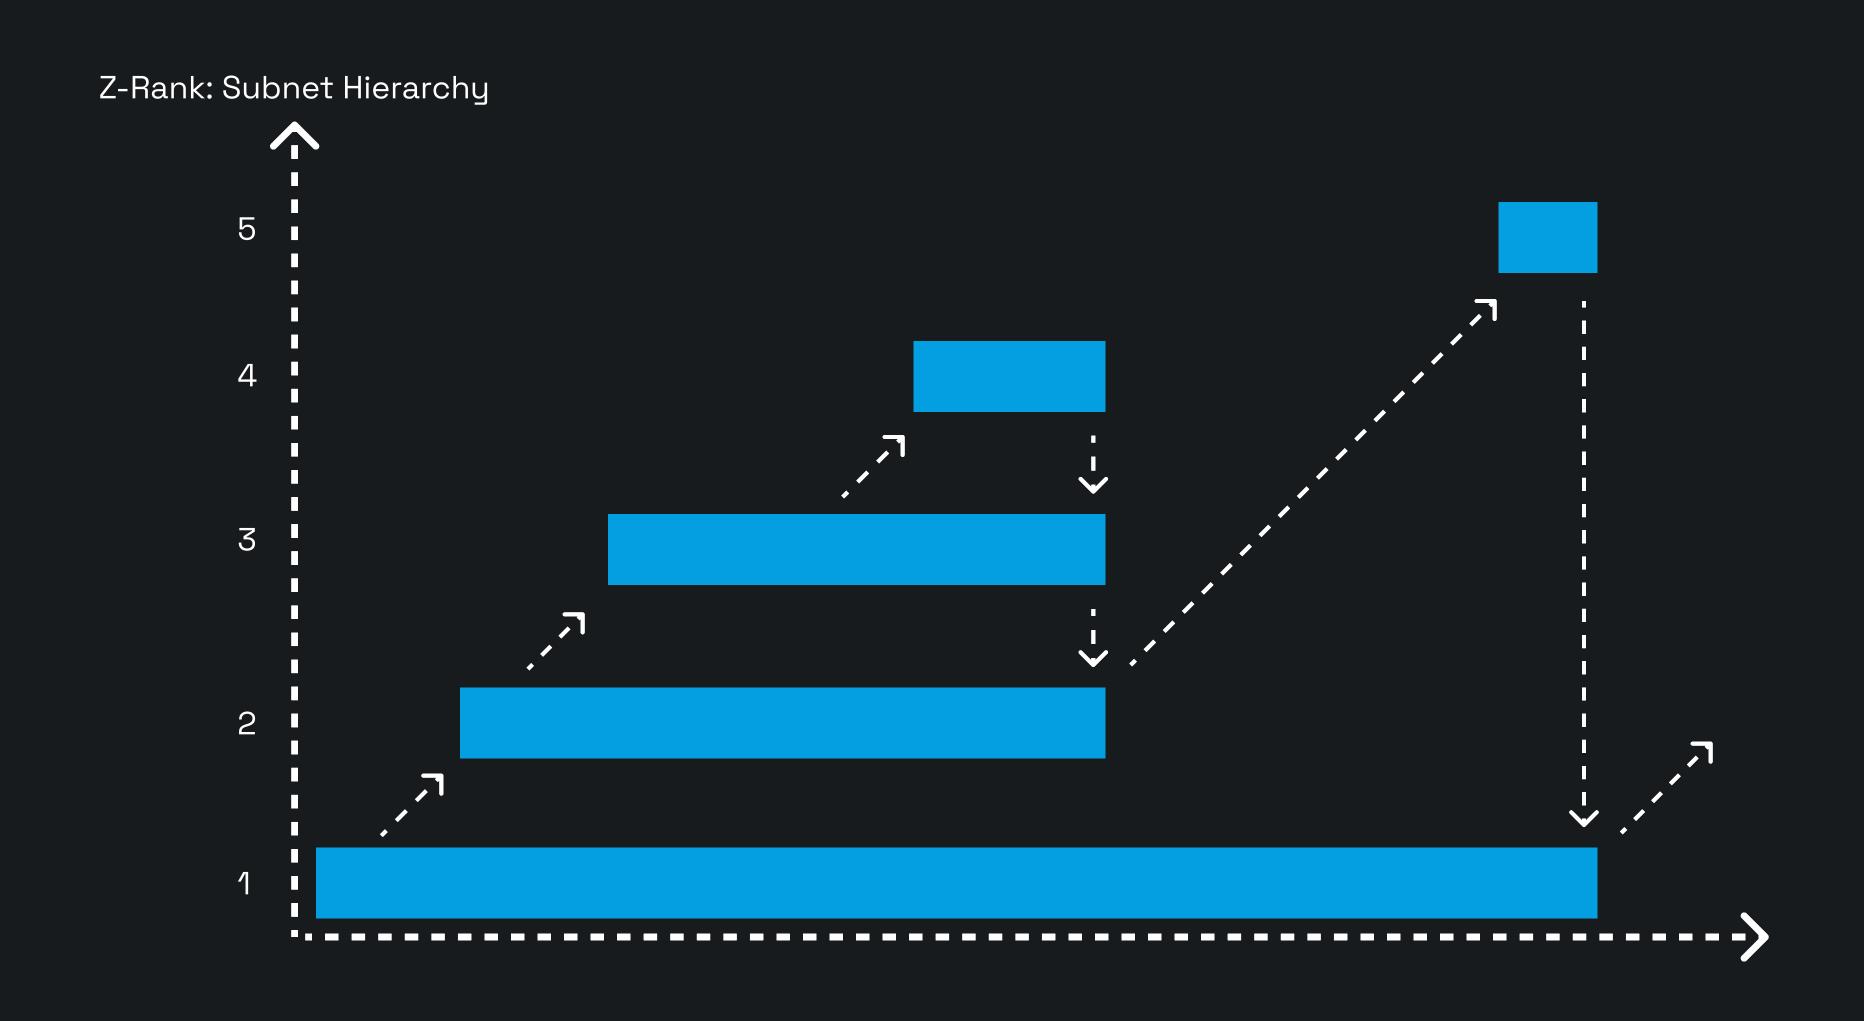


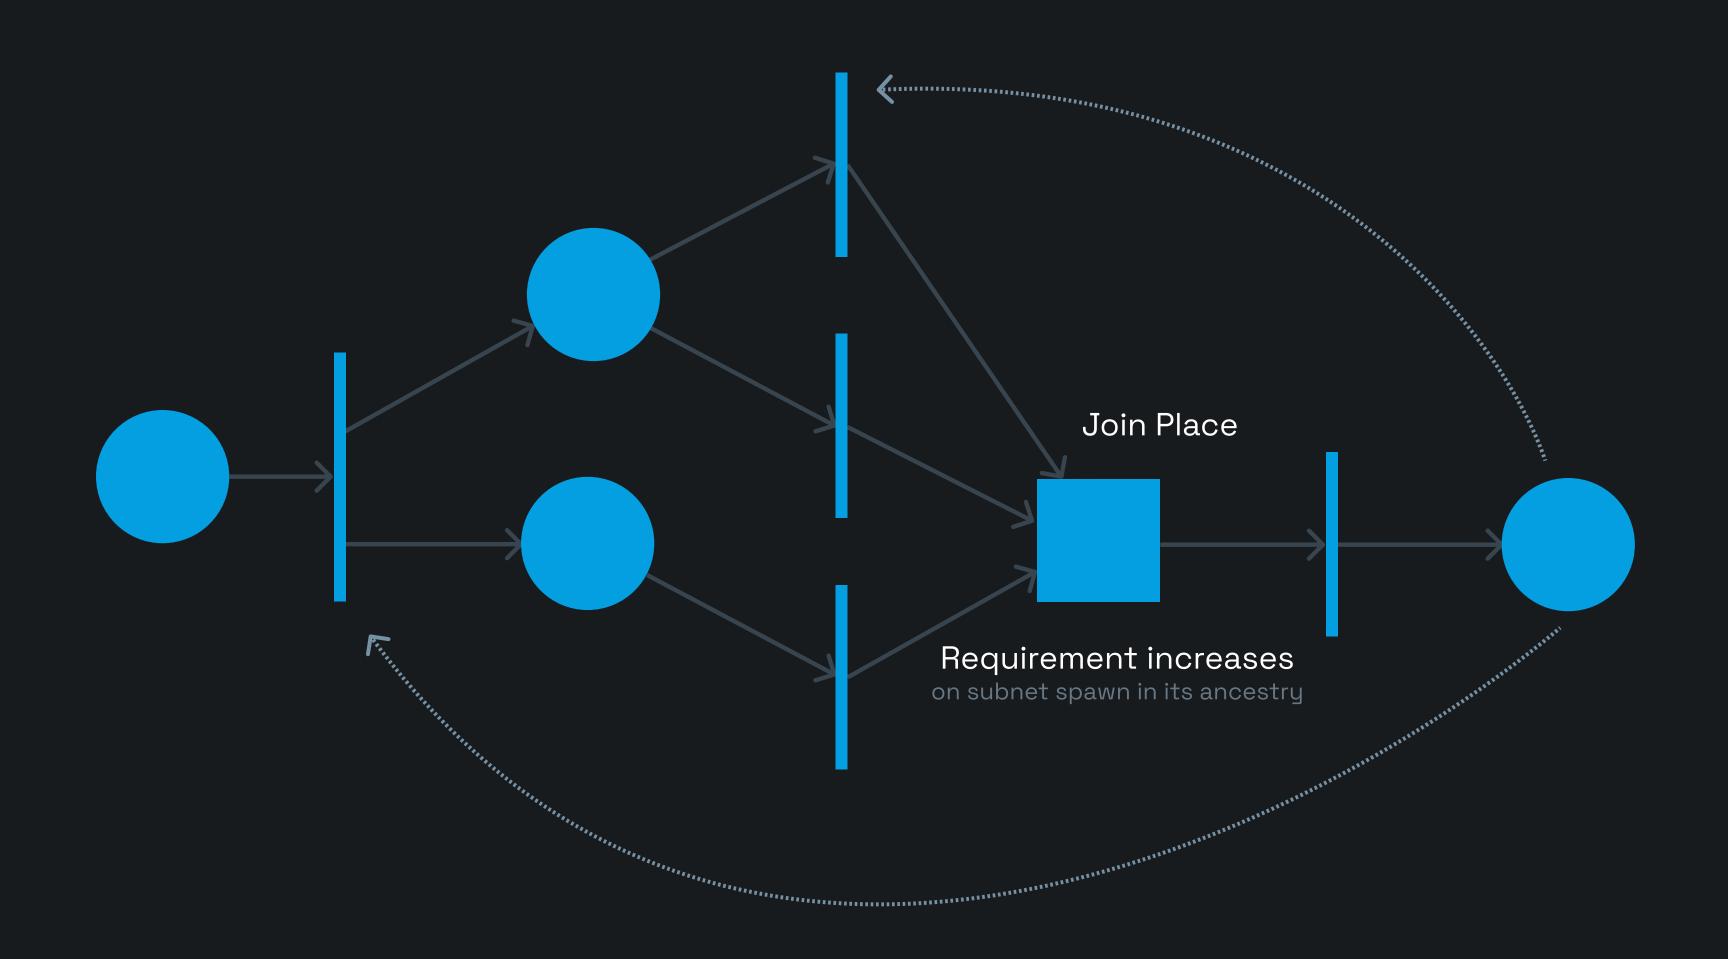
each spawned net (there can be infinite), is from a finite set of unique subnets

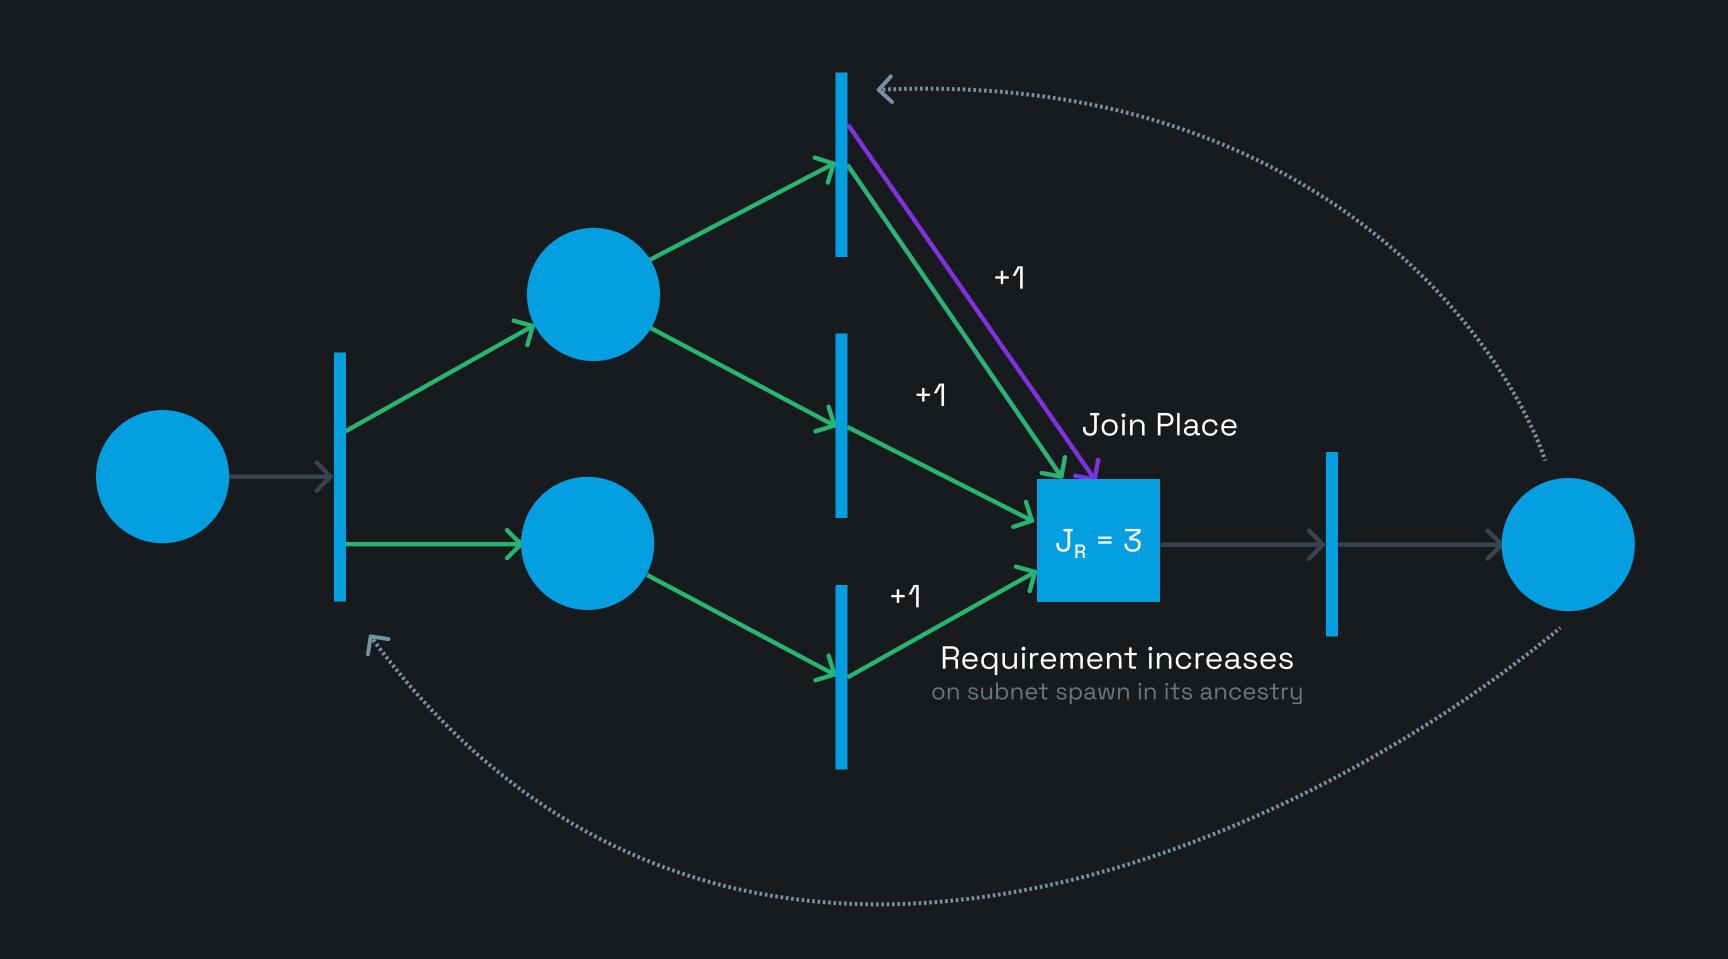


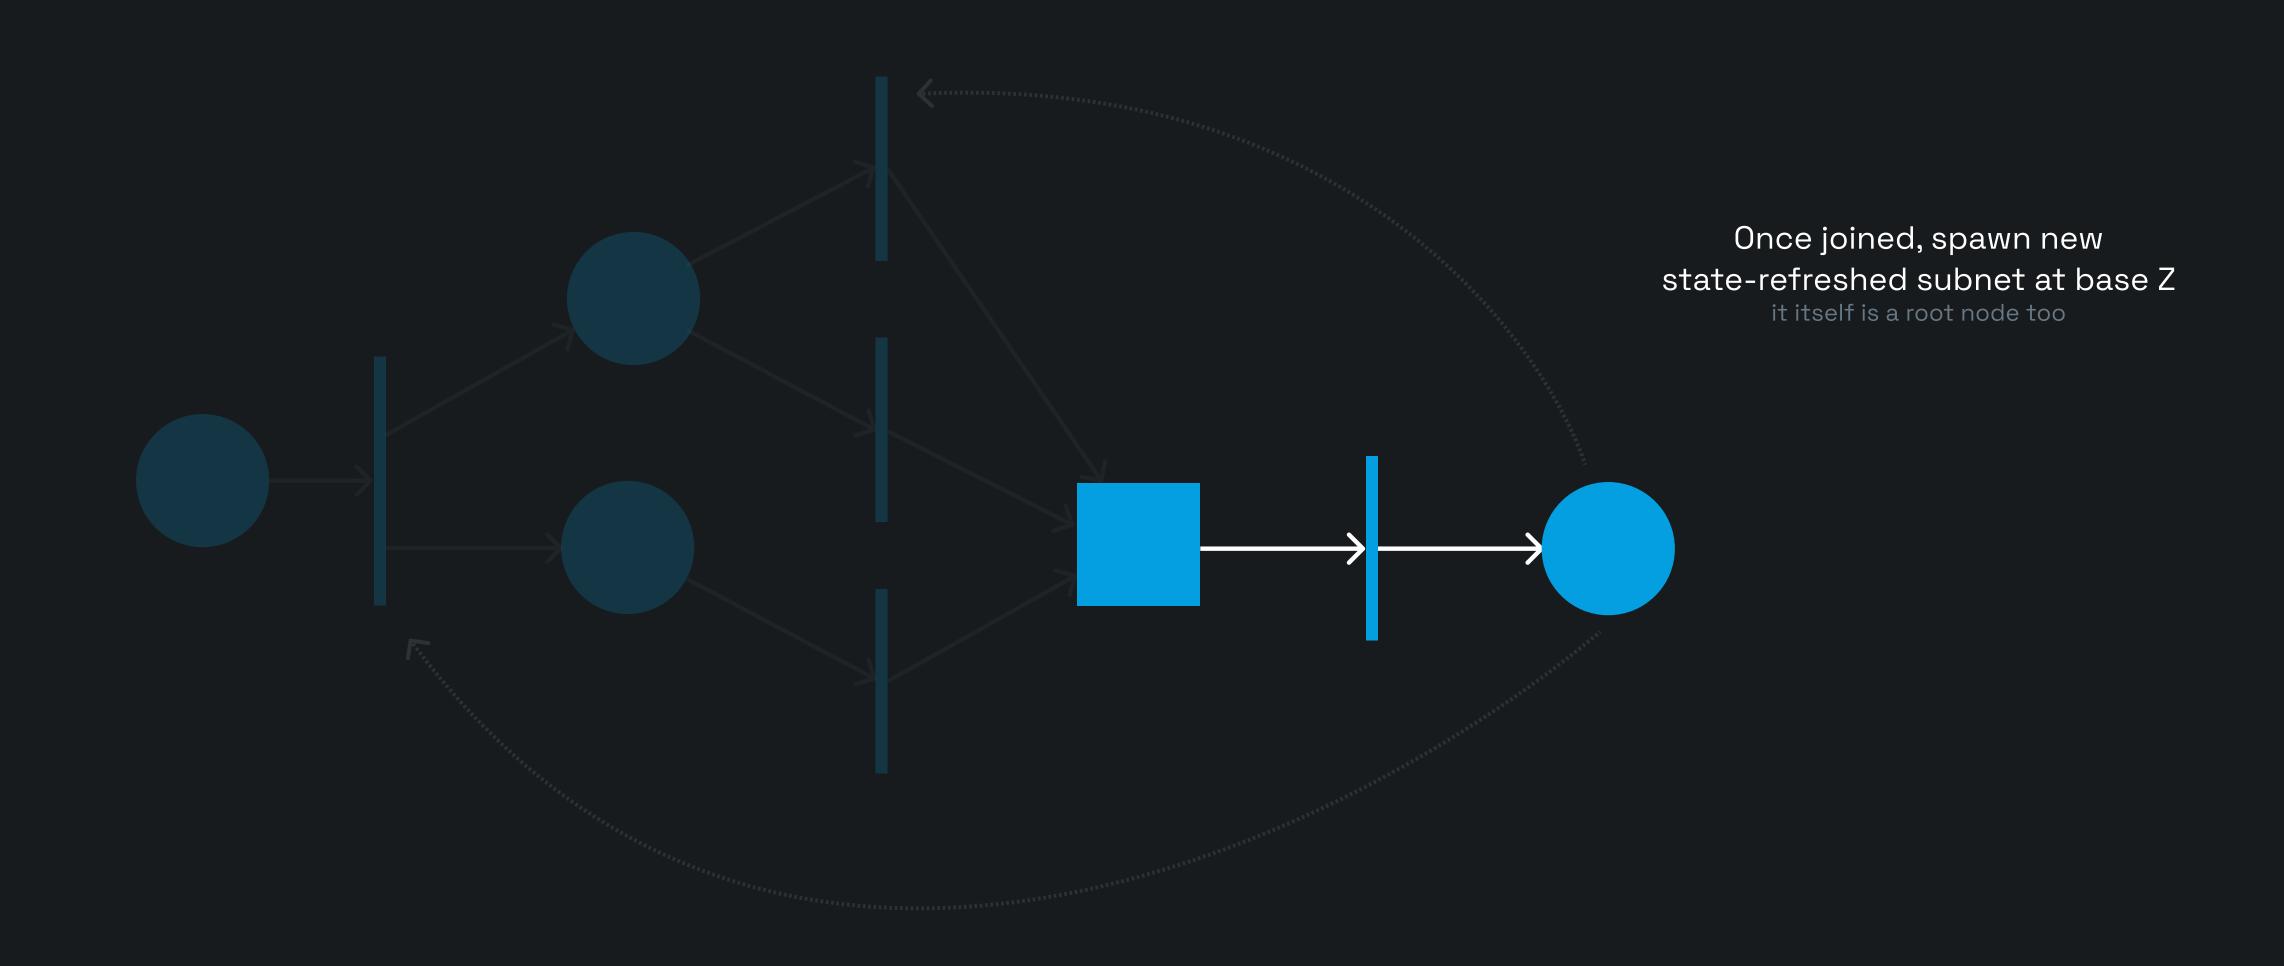


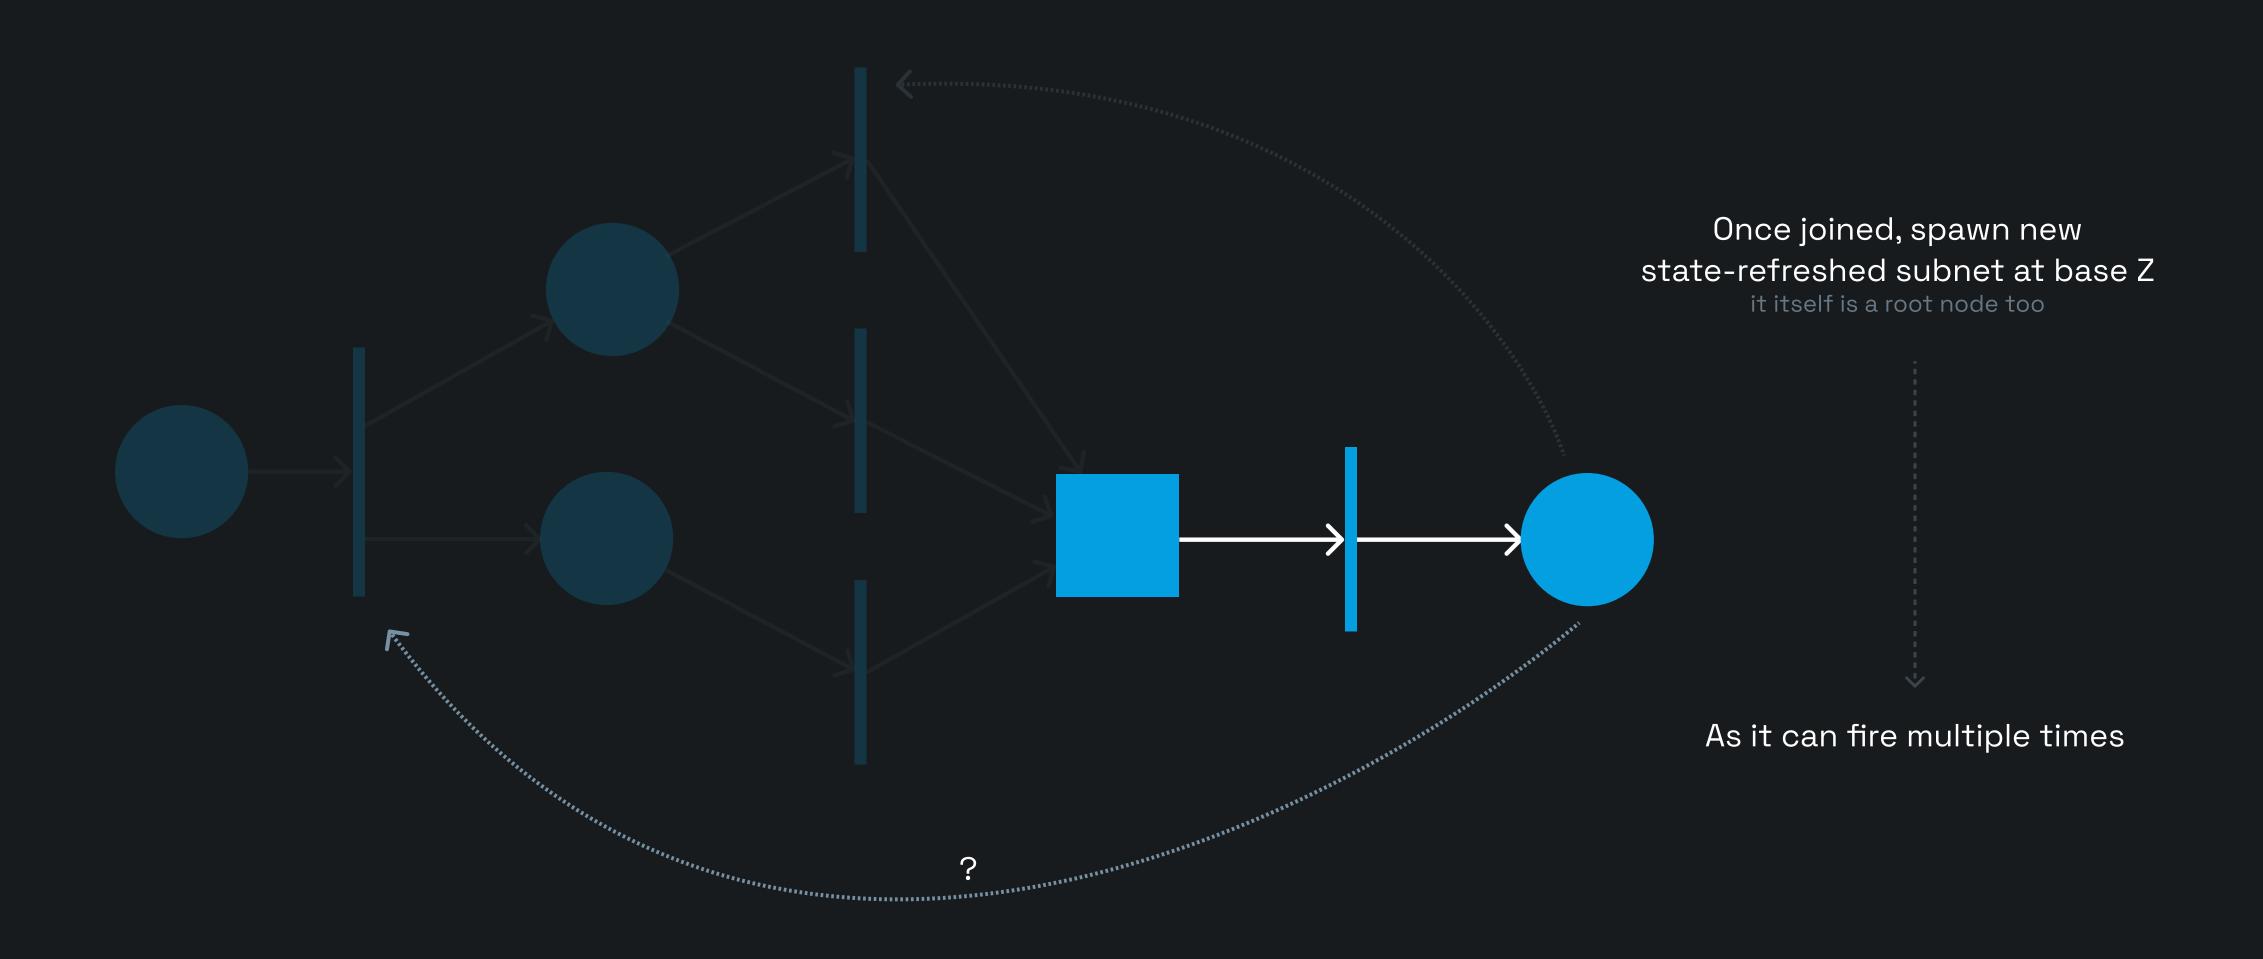




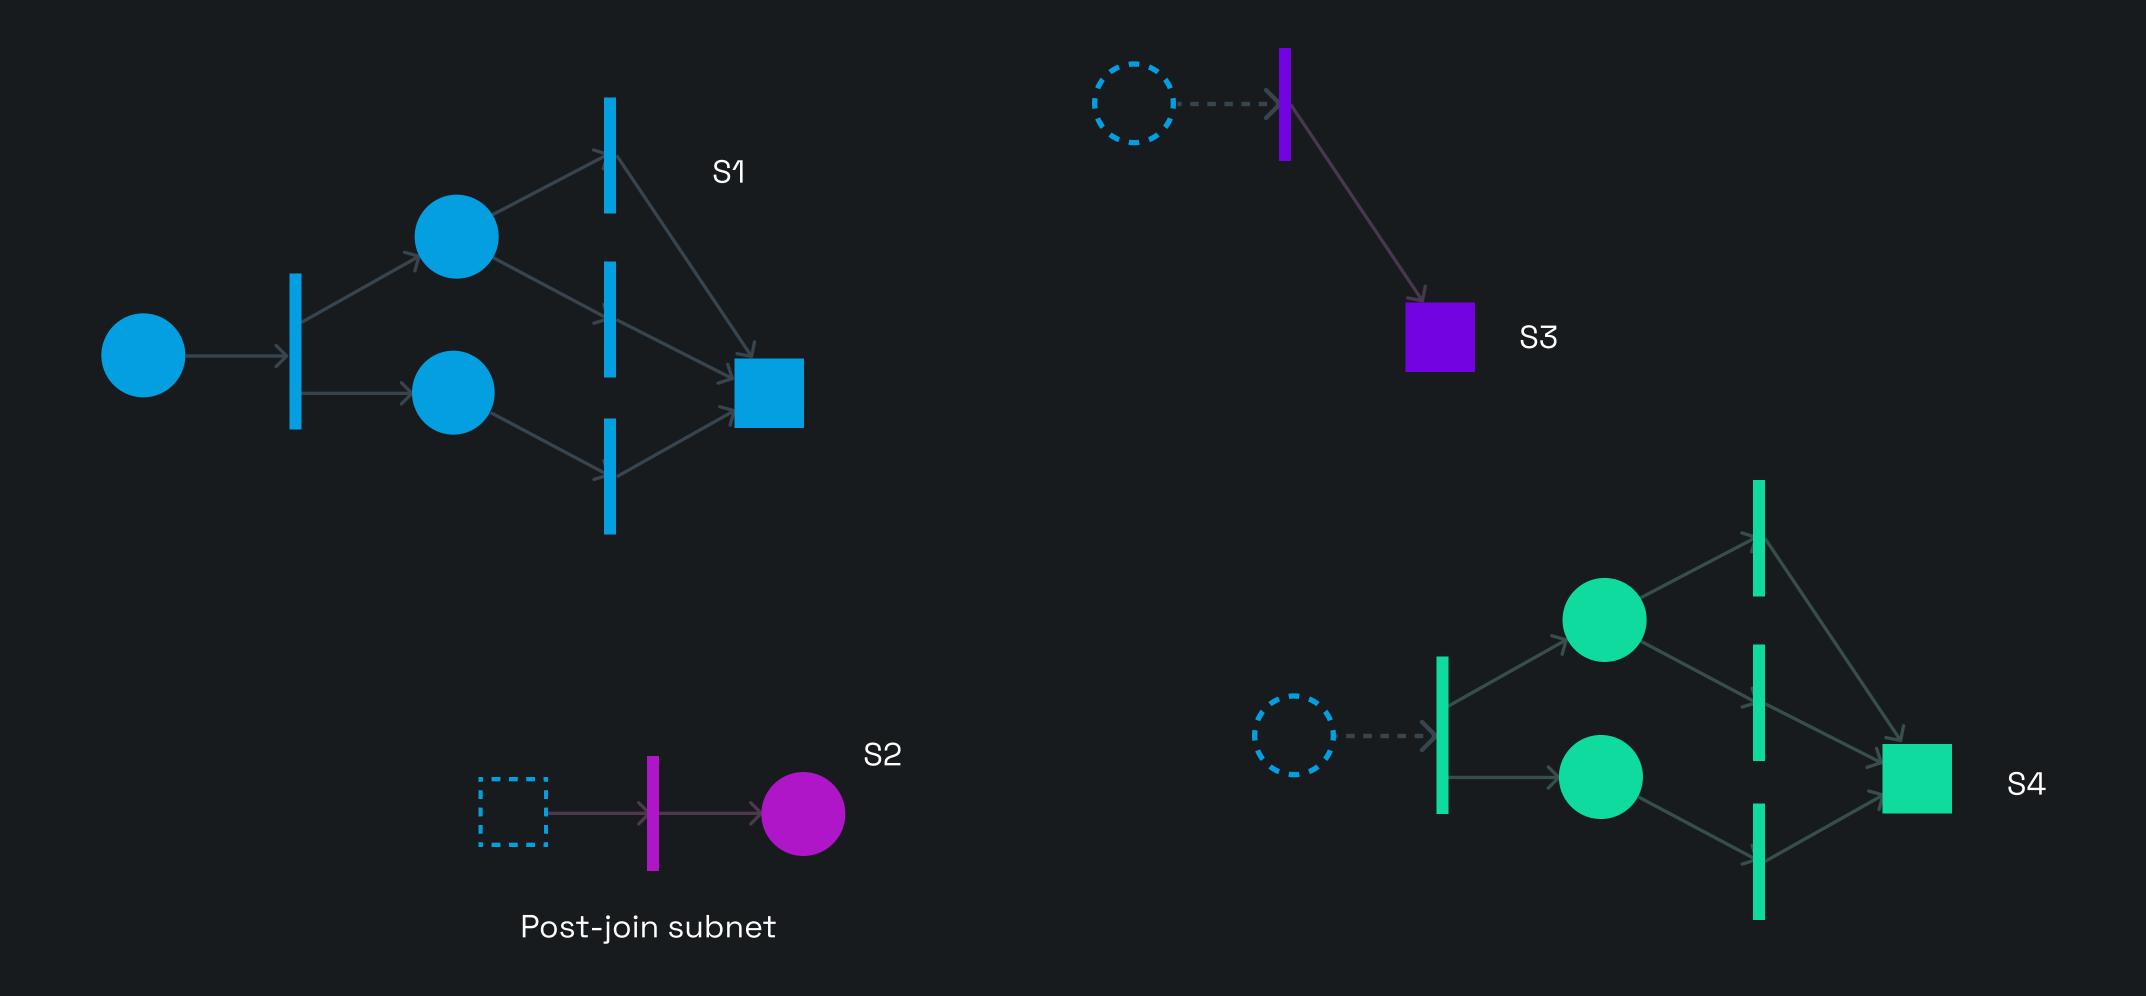




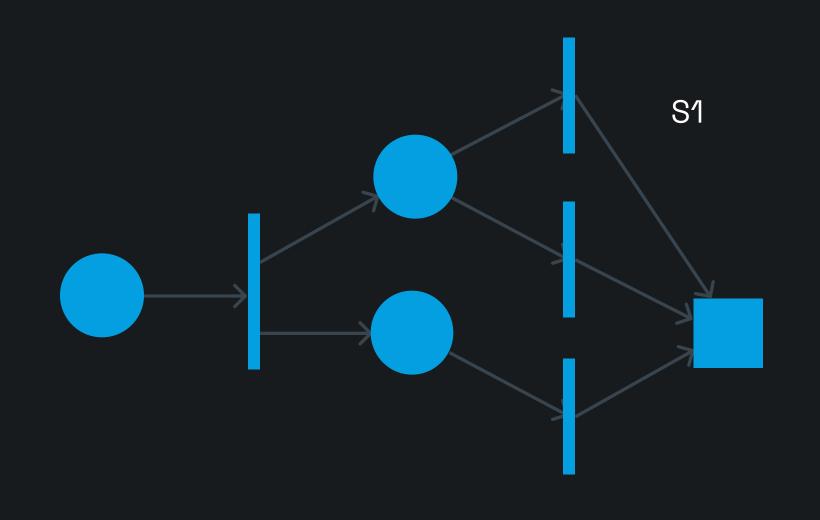


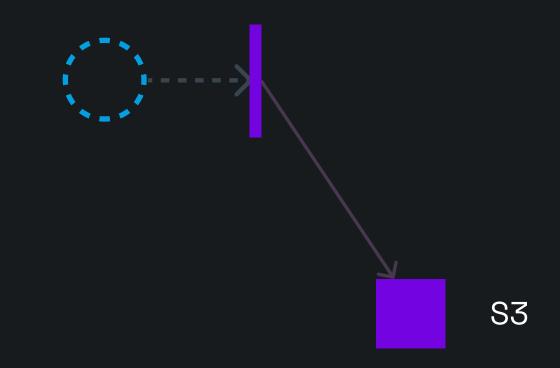


reachability for every unique acyclic subnet independently in polynomial time

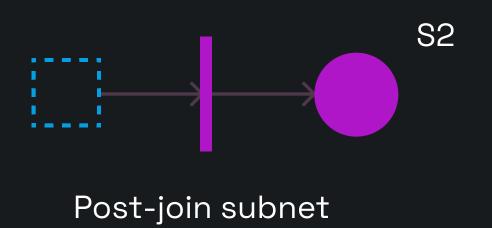


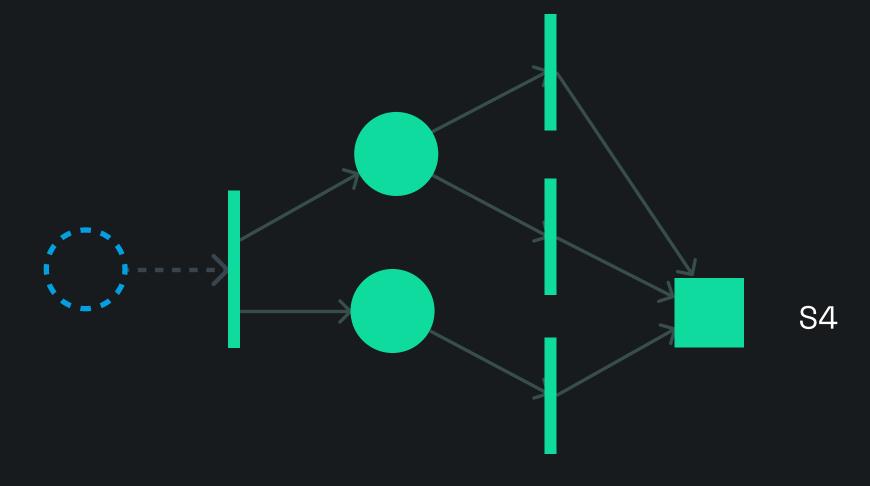
reachability for every unique acyclic subnet independently in polynomial time





 $N \in \{S1, S2, S3, S4\}$





Z-Net

 $N_z = (P, J, T, A, B, W, M_0)$

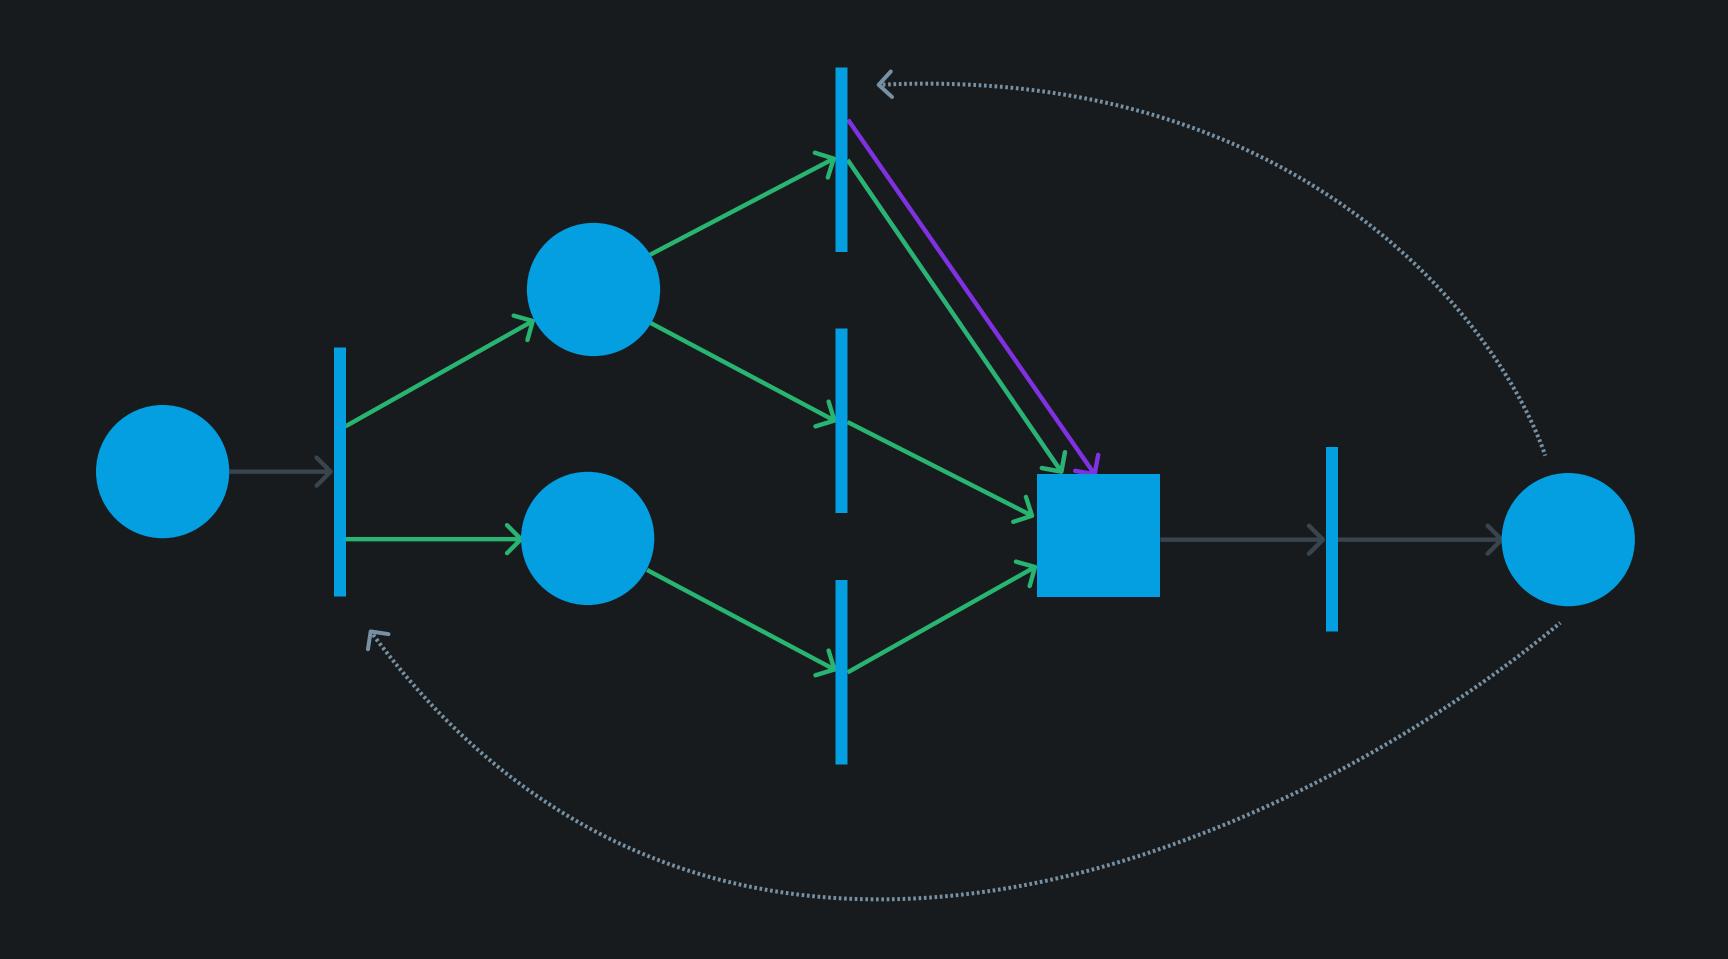
Vertices

Edges

compile-time magic

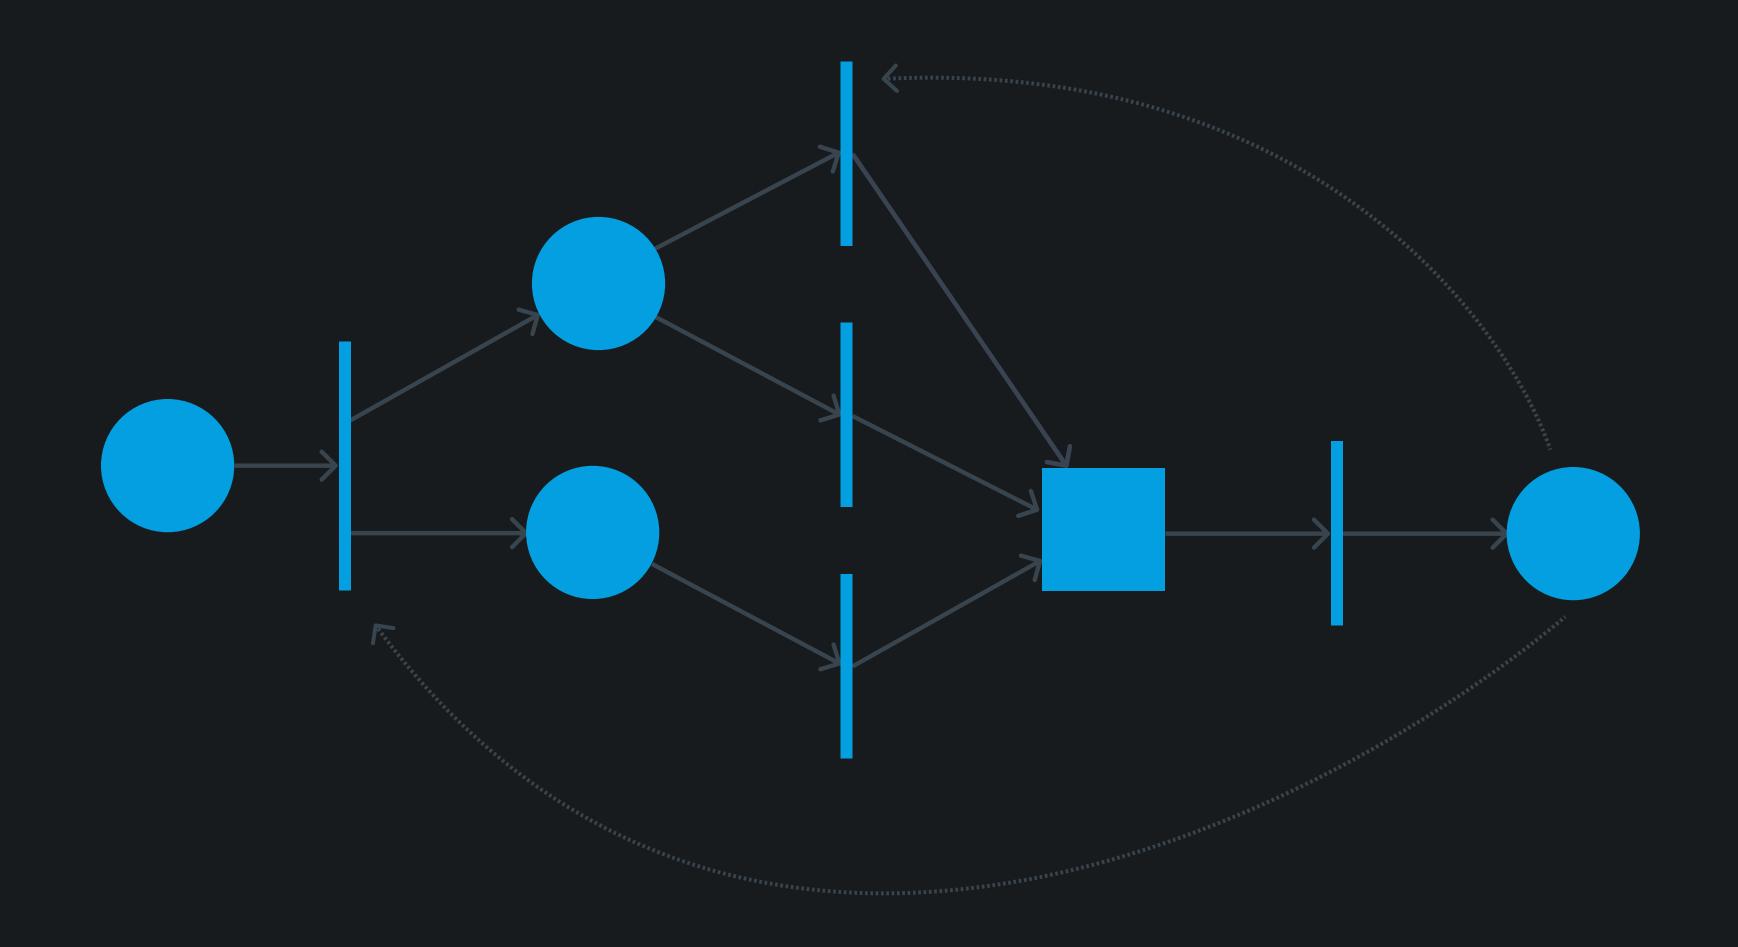
Step 1: join path resolution

traversing backwards from join node and marking each arc



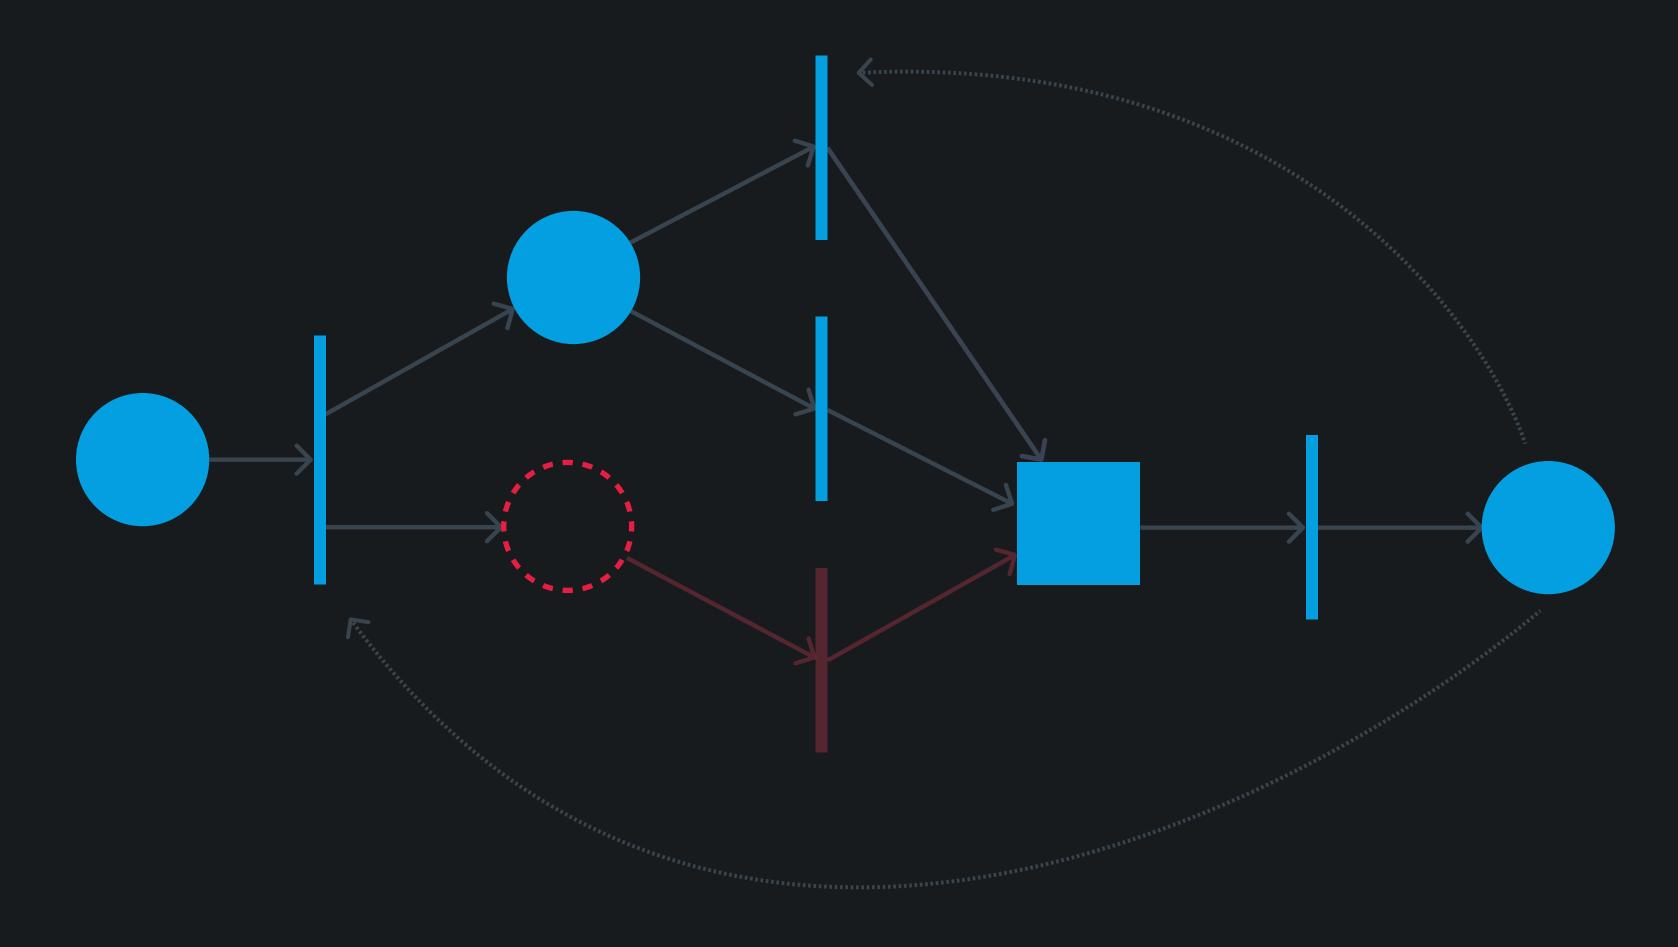
Step 2: subnet deadlock check

sort topologically, then verify all transitions are enabled by input places + check backward arcs



Extra step? Centrality measures for fault tolerance

benefits of acyclic petri nets can go beyond reachability



but a bit beyond this presentation's scope

A bit of C++

```
using System = Net<</pre>
  Places<...>,
  JoinPlaces<...>,
  Transitions<...>,
  Arcs<...>,
  BackwardArcs<...>
>;
static_assert(System::is_deadlock_free);
System system;
system.bind(T0{}, foo);
system.bind(T1{}, bar);
system.start();
```

A bit of C++

```
using System = Net<</pre>
  Places<...>,
  JoinPlaces<...>,
  Transitions<...>,
  Arcs<...>,
  BackwardArcs<...>
static_assert(System::is_deadlock_free);
System system;
system.bind(TO{}, foo);
                             ← Probably a task queue
system.bind(T1{}, bar);
                                 for transitions
system.start();
```

The end

POC and slides here: https://github.com/bdasgupta02/znet (the runtime part is in progress)